

# Dimensionality reduction

Machine Learning for Process Engineers Workshop

*Stellenbosch University*

March 2022

# Recap

- Goal of machine learning:  
*provide accurate predictions on new data*
- Model performance is a trade-off between bias and variance
- Bias can be introduced systematically through regularisation  
(sacrifice training error for testing error)
- Alternative approach to introducing bias: feature extraction

# New system

- Linear state space system: inputs  $u_1(t)$ ,  $u_2(t)$ , and measurement  $y(t)$

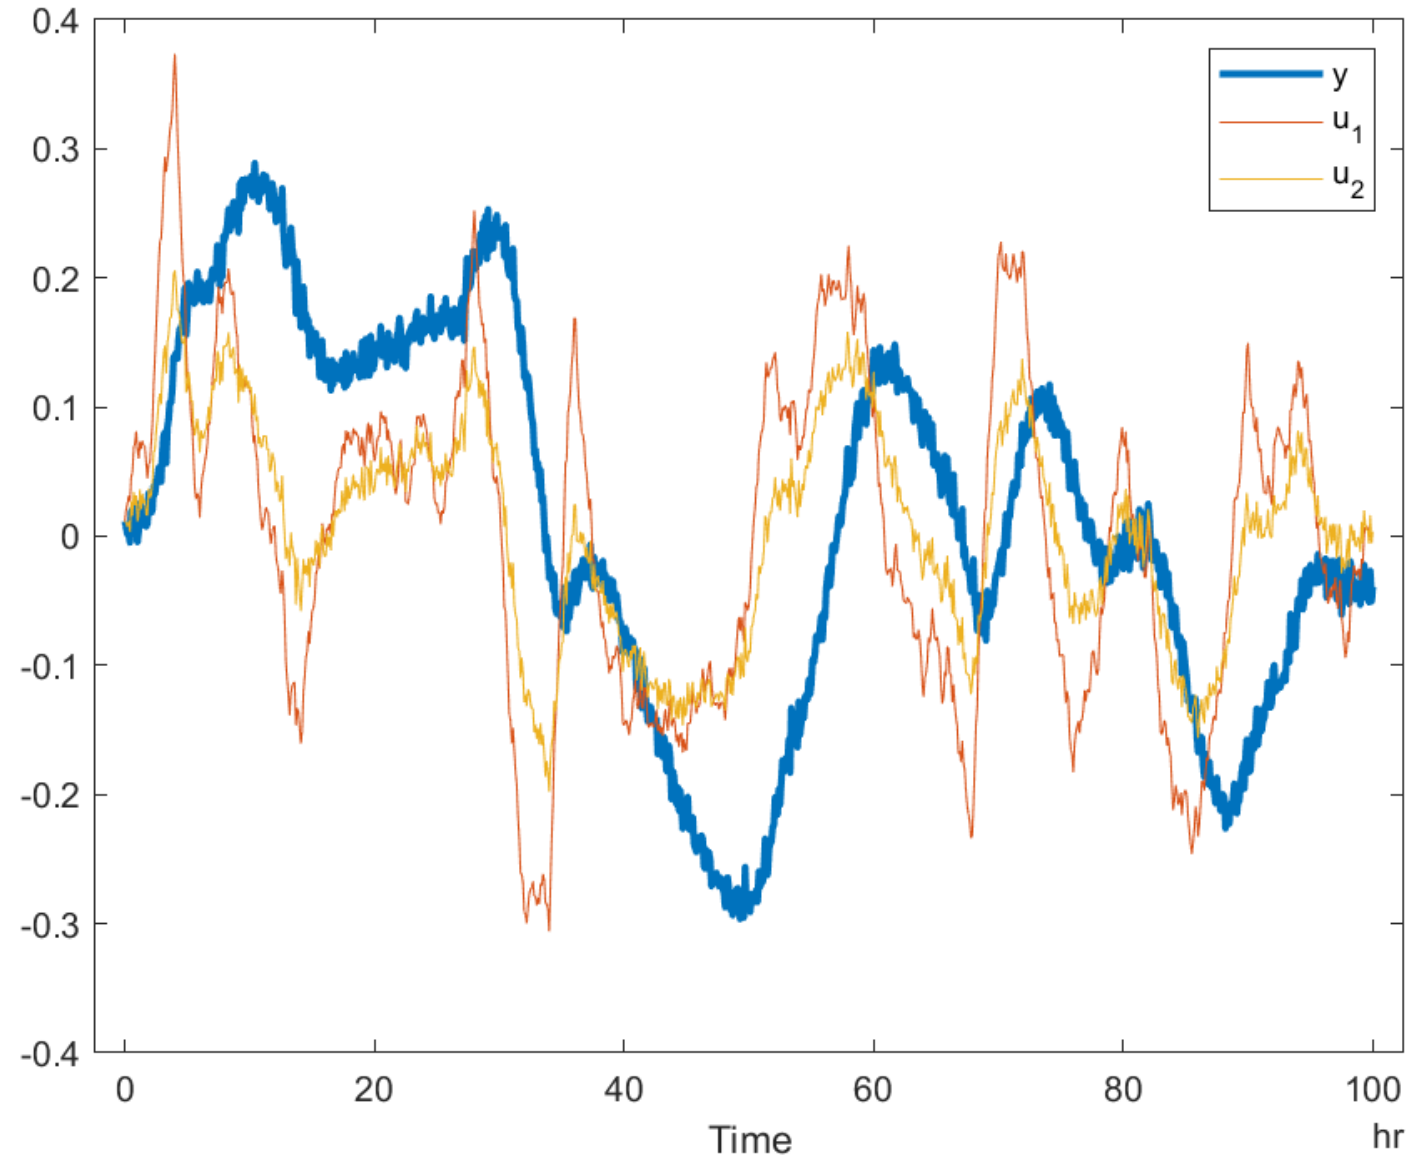
$$\frac{d\mathbf{x}}{dt} = A\mathbf{x} + B\mathbf{u}, \quad \frac{dy}{dt} = H\mathbf{x}$$

- The parameters are unknown: systems identification problem
- Simplified discrete form:

$$\begin{aligned} \hat{y}(k+1) = & \beta_1 y(k) + \beta_2 y(k-1) \dots + \beta_L y(k-L) \\ & + \beta_{L+1} u_1(k) + \beta_{L+2} u_1(k-1) \dots + \beta_{2L} u_1(k-L) \\ & + \beta_{2L+1} u_2(k) + \beta_{2L+2} u_2(k-1) \dots + \beta_{3L} u_2(k-L) \end{aligned}$$

- System has  $L$  “lags”

# New system



# New system

$$\hat{y}(L + 2) = [y(L + 1) \dots y(1) \ u_1(L + 1) \dots u_1(1) \ u_2(L + 1) \dots u_2(1)]\boldsymbol{\beta}$$

$$\hat{y}(L + 3) = [y(L + 2) \dots y(2) \ u_1(L + 2) \dots u_1(2) \ u_2(L + 2) \dots u_2(2)]\boldsymbol{\beta}$$

$$\hat{y}(L + 4) = [y(L + 3) \dots y(3) \ u_1(L + 3) \dots u_1(3) \ u_2(L + 3) \dots u_2(3)]\boldsymbol{\beta}$$

$\vdots$

$$\mathbf{y}_{L+2:N} = \mathbf{X}\boldsymbol{\beta}$$

# New system

We created a special MATLAB/Python function

```
[X, y] = CreateLaggedDesignMatrix(Data, L, f)
```

to create **y** and **X**

given time data `Data`, number of lags `L` and the fraction of data to use `f`

$$\mathbf{y}_{L+2:N} = \mathbf{X}\boldsymbol{\beta}$$

# Prediction vs simulation

- “One-step ahead” prediction problem:

$$\begin{aligned}\hat{y}(k+1) = & \beta_1 y(k) + \beta_2 y(k-1) \dots + \beta_L y(k-L) \\ & + \beta_{L+1} u_1(k) + \beta_{L+2} u_1(k-1) \dots + \beta_{2L} u_1(k-L) \\ & + \beta_{2L+1} u_2(k) + \beta_{2L+2} u_2(k-1) \dots + \beta_{3L} u_2(k-L)\end{aligned}$$

- Simulation problem

$$\begin{aligned}\hat{y}(k+1) = & \beta_1 \hat{y}(k) + \beta_2 \hat{y}(k-1) \dots + \beta_L \hat{y}(k-L) \\ & + \beta_{L+1} u_1(k) + \beta_{L+2} u_1(k-1) \dots + \beta_{2L} u_1(k-L) \\ & + \beta_{2L+1} u_2(k) + \beta_{2L+2} u_2(k-1) \dots + \beta_{3L} u_2(k-L)\end{aligned}$$

- Simulation is more difficult

# Prediction vs simulation

We created a special MATLAB/Python function

```
y = PredictTimeSeries mdl, Data, L)
```

to generate predictions  $y$

given a model  $mdl$ , time data  $Data$ , and number of lags  $L$

The model consists of  $mdl.beta$  containing the vector of coefficients  $\beta$   
and  $mdl.Q$  containing a transformation matrix  $Q$  (to be discussed)



# New system

- Simplified discrete form:

$$\begin{aligned}\hat{\mathbf{y}}(\mathbf{k} + \mathbf{1}) = & \beta_1 \hat{\mathbf{y}}(\mathbf{k}) + \beta_2 \hat{\mathbf{y}}(\mathbf{k} - \mathbf{1}) \dots + \beta_L \hat{\mathbf{y}}(\mathbf{k} - L) \\ & + \beta_{L+1} u_1(k) + \beta_{L+2} u_1(k - 1) \dots + \beta_{2L} u_1(k - L) \\ & + \beta_{2L+1} u_2(k) + \beta_{2L+2} u_2(k - 1) \dots + \beta_{3L} u_2(k - L)\end{aligned}$$

- System above has  $L$  “lags”  
(often used when modelling time series data)
- Not sure how many lags are required
- Consecutive data points are highly correlated

# In MATLAB/Python: Example 10

- Open file “MLforProcEng\_Workshop\_3.m” and run the  
`%% Initialize`
- Go to the cell  
`%% Example 10: Linear model fit to  
timeseries data`
- The code is ready to run and needs no adjustment
- Note the construction of the “mdl” object for use in  
“PredictTimeSeries”  
`mdl.Q = 1` in both cases (for now)

```
%% Initialize
:
load ProcessData

% Prepare the design matrix "X"
L = 10;
:
[X, y] = CreateLaggedDesignMatrix(Data, L, 0.1);

%% Example 10: Linear model fit to timeseries data
% Fit a linear model without regularization
mdl = fitlm(X, y, 'Intercept', false);
linear_mdl.Q = 1;
linear_mdl.beta = mdl.Coefficients.Estimate;
y_linear = PredictTimeSeries(linear_mdl, Data, L);

% Fit a linear model with ridge regression
?
?
?
:
```

```
%% Initialize
:
load ProcessData

% Prepare the design matrix "X"
L = 10;
:
[X, y] = CreateLaggedDesignMatrix(Data, L, 0.1);

%% Example 10: Linear model fit to timeseries data
% Fit a linear model without regularization
mdl = fitlm(X, y, 'Intercept', false);
linear_mdl.Q = 1;
linear_mdl.beta = mdl.Coefficients.Estimate;
y_linear = PredictTimeSeries(linear_mdl, Data, L);

% Fit a linear model with ridge regression
ridge_mdl.Q = 1;
ridge_mdl.beta = lasso(X, y, 'Alpha', 1e-6, 'Lambda', 0.1);
y_ridge = PredictTimeSeries(ridge_mdl, Data, L);
:
```

```

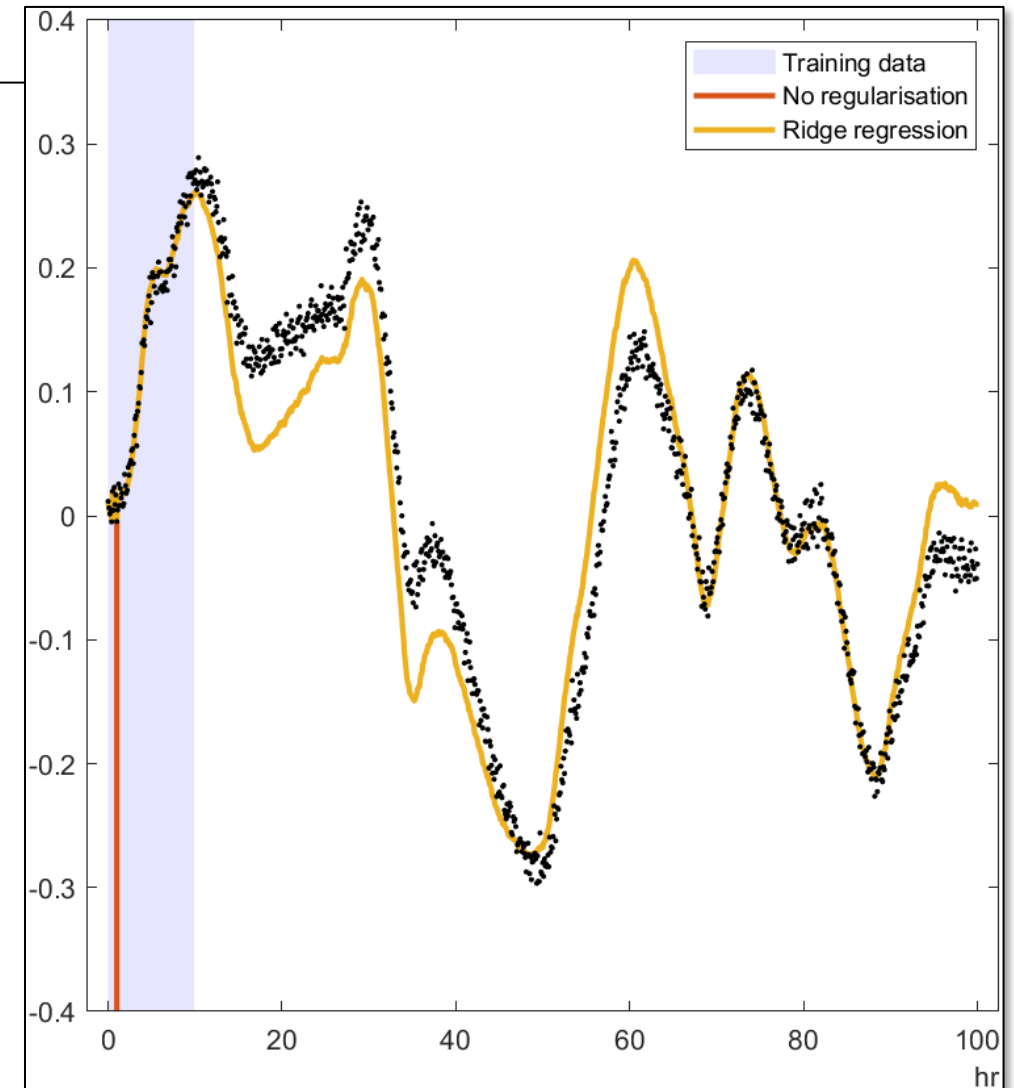
%% Initialize
:
load ProcessData

% Prepare the design matrix "X"
L = 10;
:
[X, y] = CreateLaggedDesignMatrix(Data, L, 0.1);

%% Example 10: Linear model fit to timeseries data
% Fit a linear model without regularization
mdl = fitlm(X, y, 'Intercept', false);
linear_mdl.Q = 1;
linear_mdl.beta = mdl.Coefficients.Estimate;
y_linear = PredictTimeSeries(linear_mdl, Data, L);

% Fit a linear model with ridge regression
ridge_mdl.Q = 1;
ridge_mdl.beta = lasso(X, y, 'Alpha', 1e-6, 'Lambda', 0.1);
y_ridge = PredictTimeSeries(ridge_mdl, Data, L);
:

```



# Correlation in predictors

- The linear regression problem solves the following:

$$\mathbf{X}^T \mathbf{y} = \mathbf{X}^T \mathbf{X} \boldsymbol{\beta}$$

- If two predictors are perfectly correlated (e.g.  $\mathbf{X}_2 = \gamma \mathbf{X}_1$ ) then  $\mathbf{X}^T \mathbf{X}$  is singular
- If two predictors are closely correlated (e.g.  $\mathbf{X}_2 \approx \gamma \mathbf{X}_1$ ) then  $\mathbf{X}^T \mathbf{X}$  is ill-conditioned

# Correlation in predictors

- The linear regression problem solves the following:

$$\mathbf{X}^T \mathbf{y} = \mathbf{X}^T \mathbf{X} \boldsymbol{\beta}$$

- If two predictors are perfectly correlated (e.g.  $\mathbf{X}_2 = \gamma \mathbf{X}_1$ ) then  $\mathbf{X}^T \mathbf{X}$  is singular
- If two predictors are closely correlated (e.g.  $\mathbf{X}_2 \approx \gamma \mathbf{X}_1$ ) then  $\mathbf{X}^T \mathbf{X}$  is ill-conditioned
- The ridge regression problem solves the following:

$$\mathbf{X}^T \mathbf{y} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}) \boldsymbol{\beta}$$

- Addition of the  $\lambda$  on the diagonal conditions the matrix

# Correlation in predictors

- The linear regression problem solves the following:

$$\mathbf{X}^T \mathbf{y} = \mathbf{X}^T \mathbf{X} \boldsymbol{\beta}$$

- If two predictors are perfectly correlated (e.g.  $\mathbf{X}_2 = \gamma \mathbf{X}_1$ ) then  $\mathbf{X}^T \mathbf{X}$  is singular
- If two predictors are closely correlated (e.g.  $\mathbf{X}_2 \approx \gamma \mathbf{X}_1$ ) then  $\mathbf{X}^T \mathbf{X}$  is ill-conditioned
- The ridge regression problem solves the following:

$$\mathbf{X}^T \mathbf{y} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}) \boldsymbol{\beta}$$

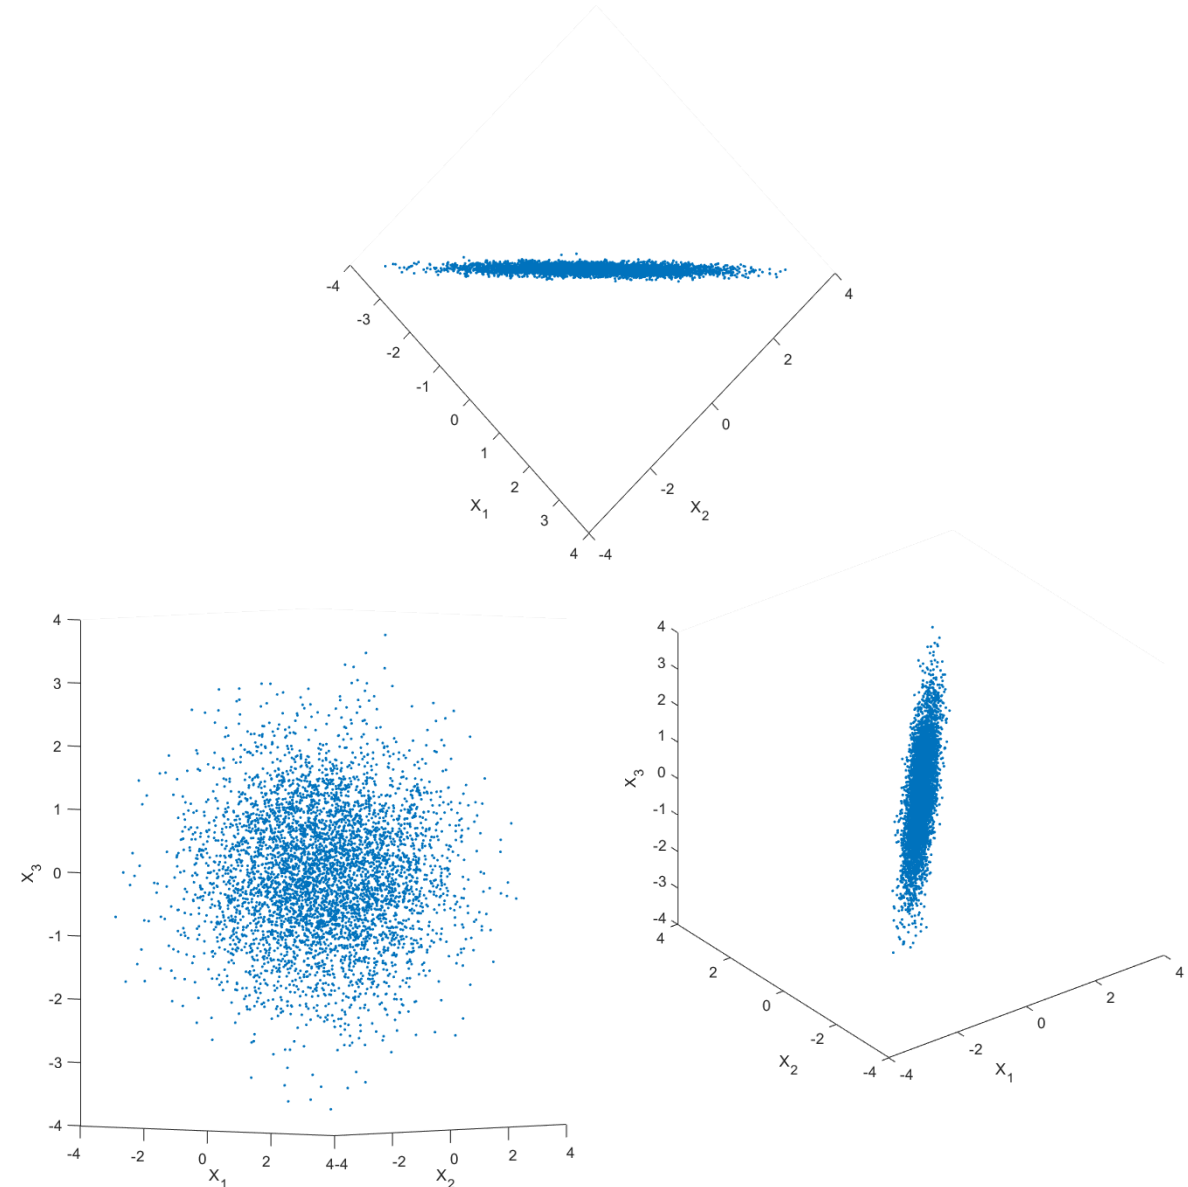
- Addition of the  $\lambda$  on the diagonal conditions the matrix

Interestingly “`lasso(X, y, 'Lambda', 0);`” yields *much* better results than “`fitlm(X, y);`” due to the underlying numerical method



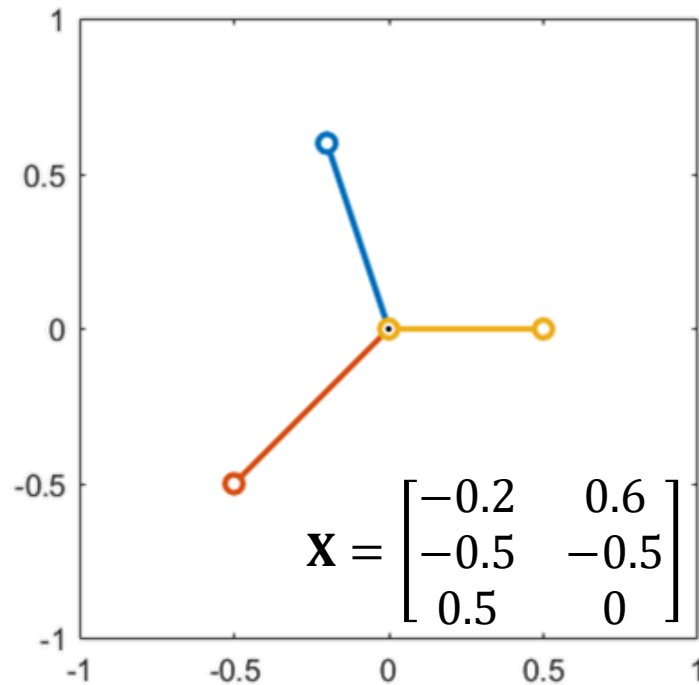
# Projection of predictors

- Can we “combine” correlated variables into single components?



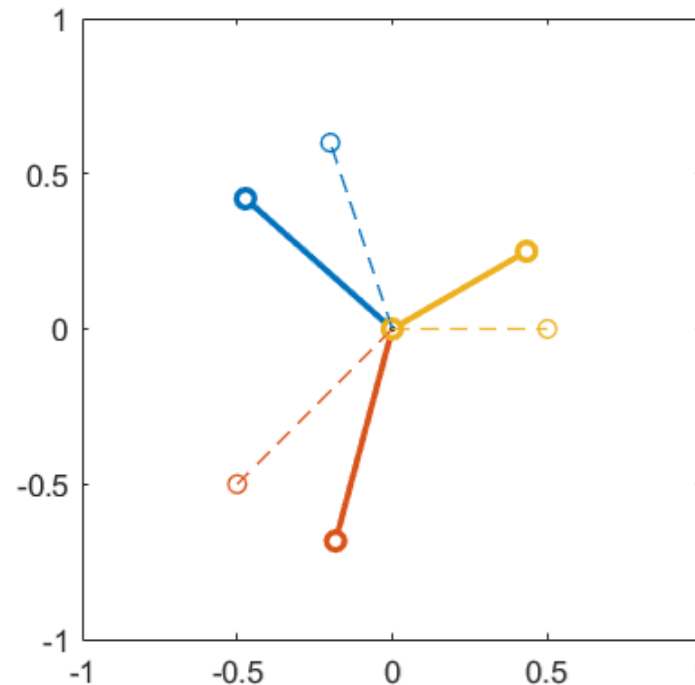
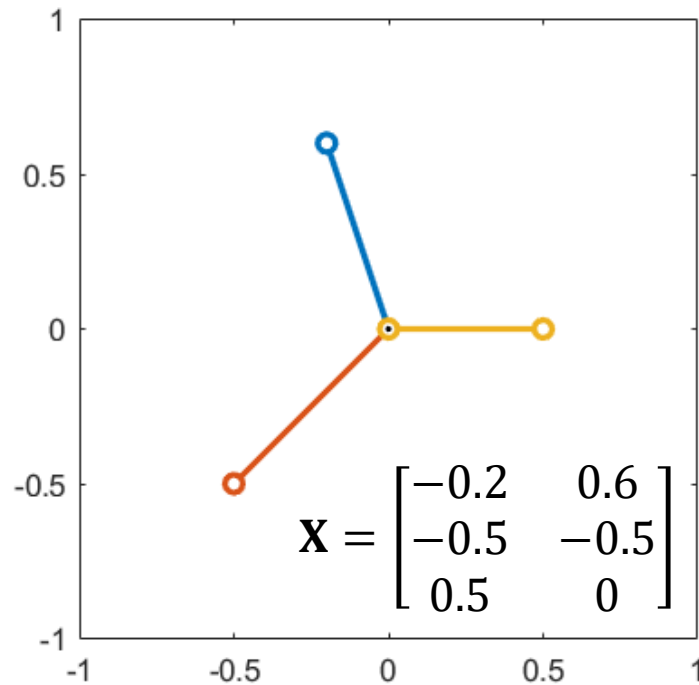
# Orthonormal matrix as “rotation”

- Matrices can project vectors to a new set of basis functions
- Orthonormal matrices amounts to a rotation of the basis vectors



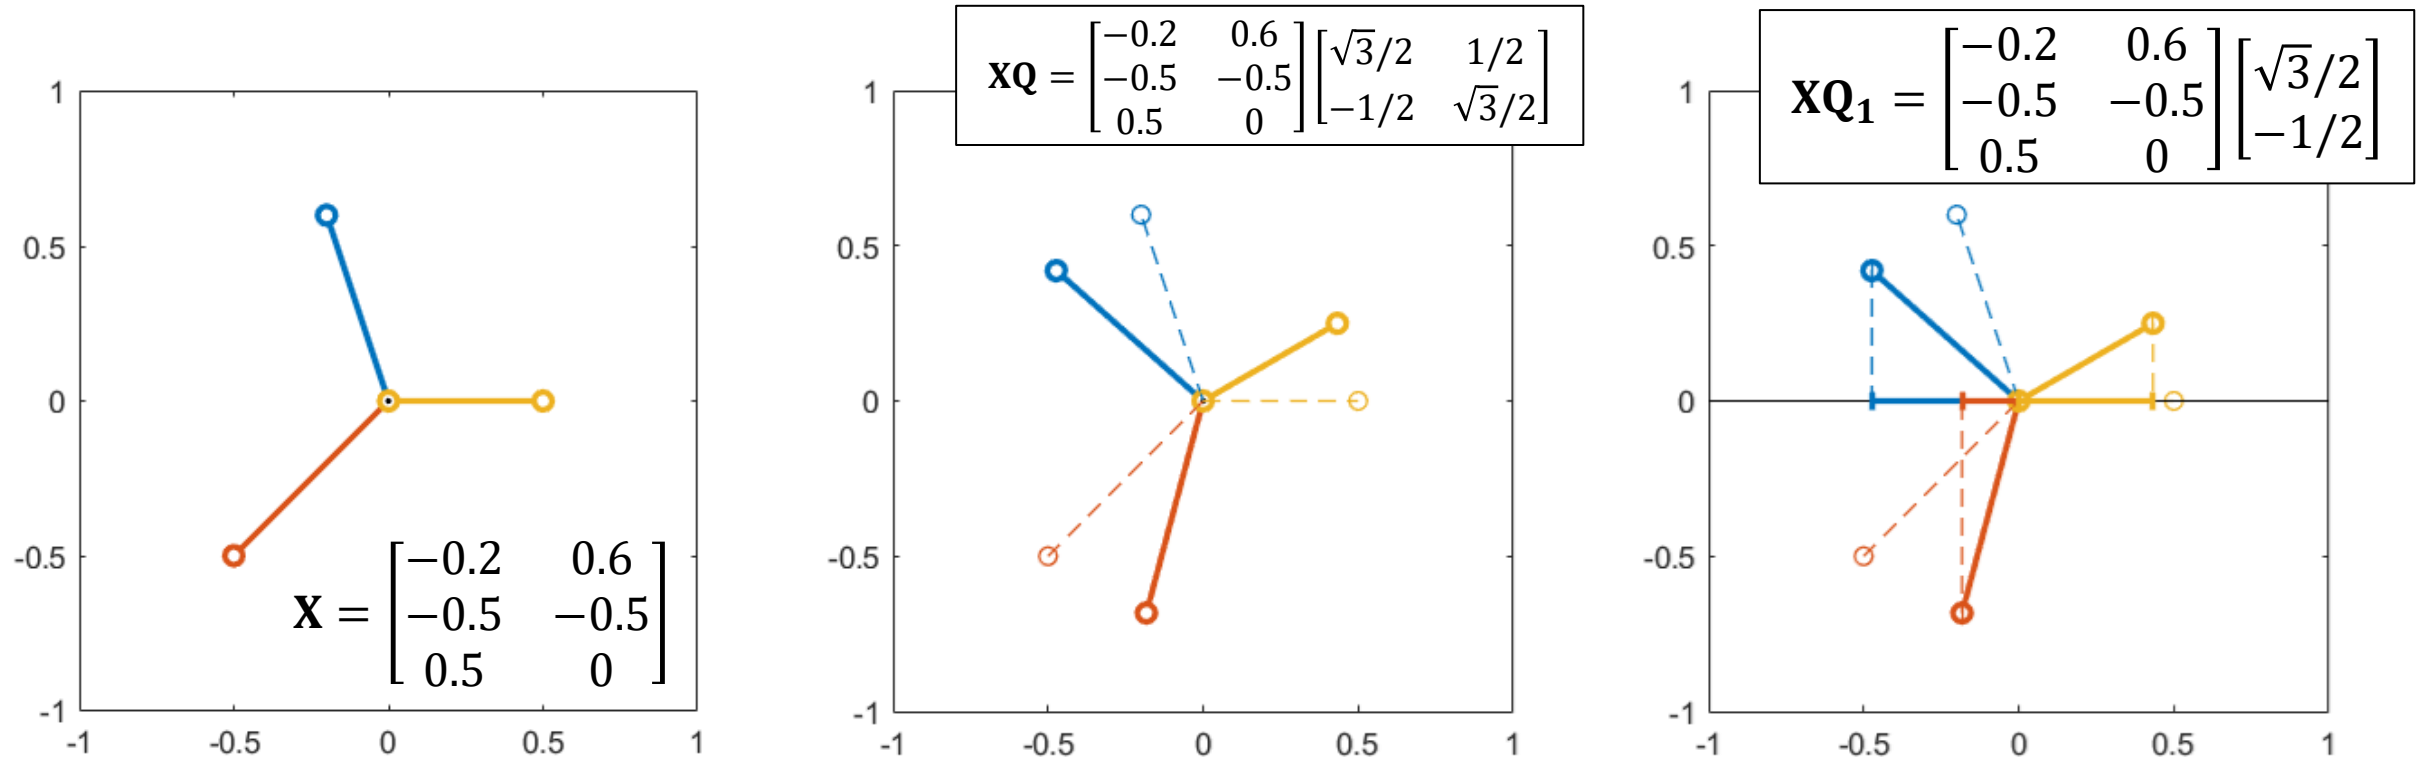
# Orthonormal matrix as “rotation”

- Matrices can project vectors to a new set of basis functions
- Orthonormal matrices amounts to a rotation of the basis vectors



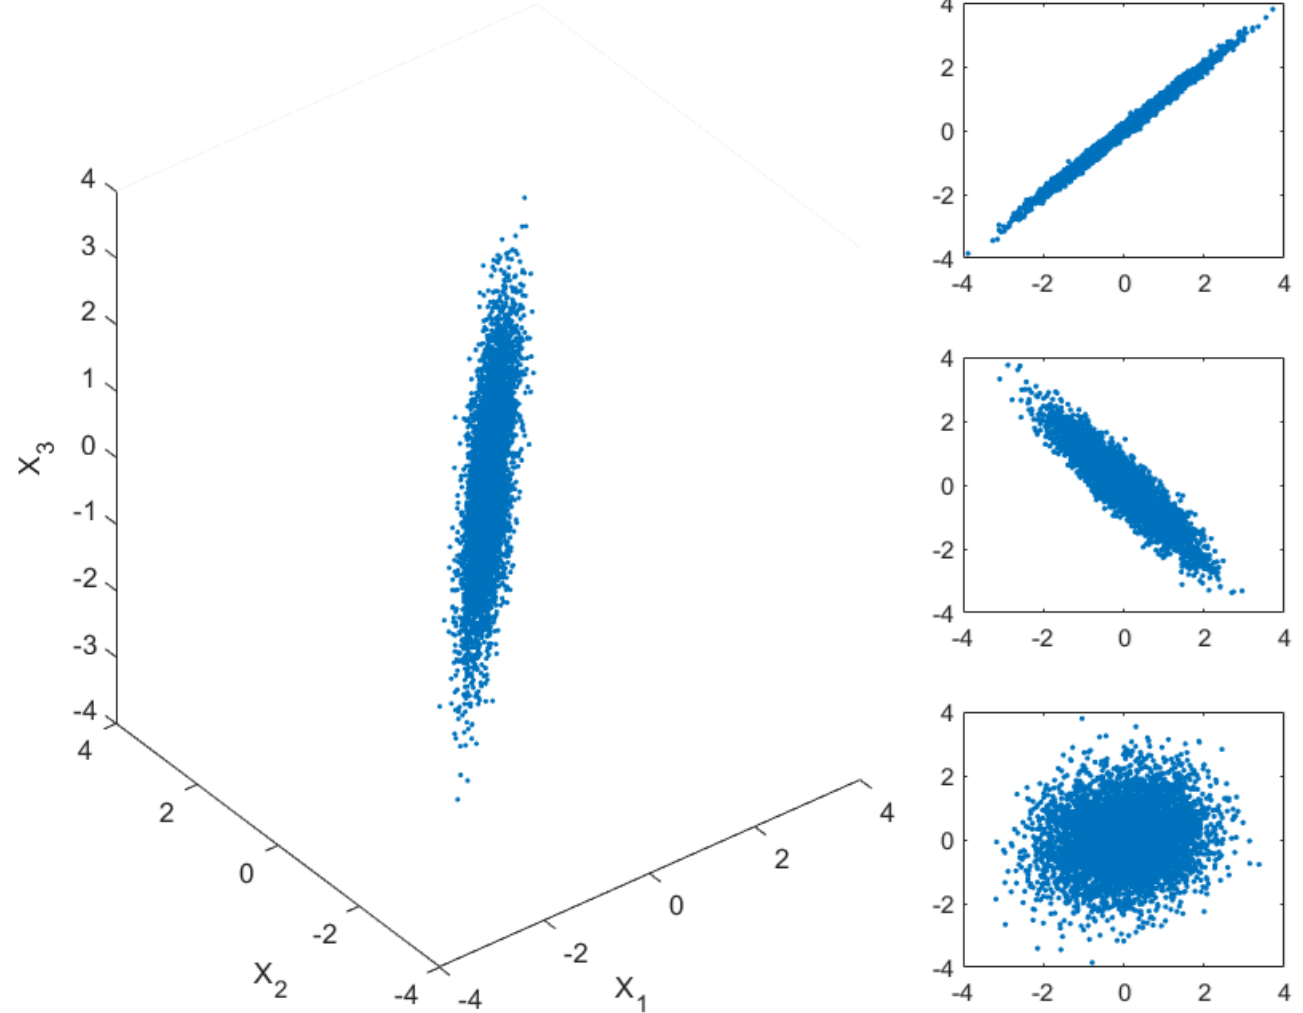
# Orthonormal matrix as “rotation”

- Matrices can project vectors to a new set of basis functions
- Orthonormal matrices amounts to a rotation of the basis vectors

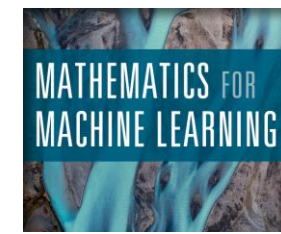


# Orthonormal matrix as “rotation”

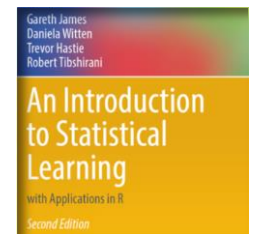
- What “projection” (rotation, eliminating dimensions) of the data yields a reduced dimension data set with the greatest *variance*?



# Orthonormal matrix as “rotation”

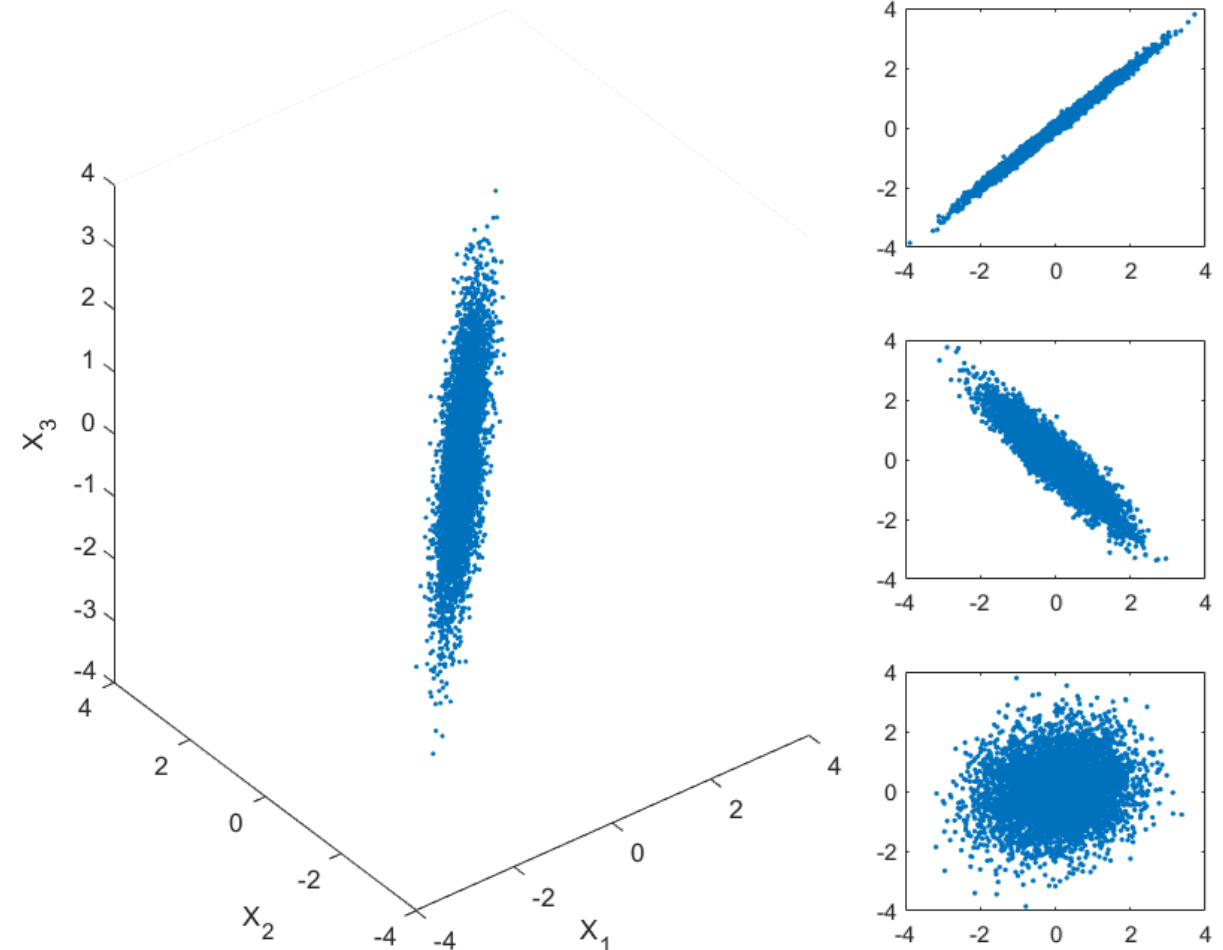


p 115

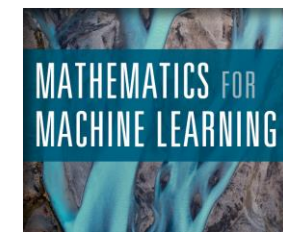


p 498

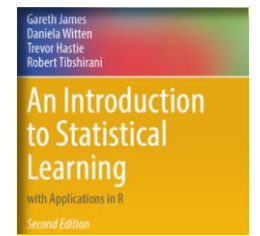
- What “projection” (rotation, eliminating dimensions) of the data yields a reduced dimension data set with the greatest *variance*?
- **Principal Component Analysis (PCA)**: Use the *eigenvectors*  $\mathbf{q}_1, \mathbf{q}_2, \mathbf{q}_3 \dots$  of the *covariance matrix*  $\mathbf{X}^T \mathbf{X}$  corresponding to the largest *eigenvalues*  $\lambda_1, \lambda_2, \lambda_3 \dots$



# Orthonormal matrix as “rotation”

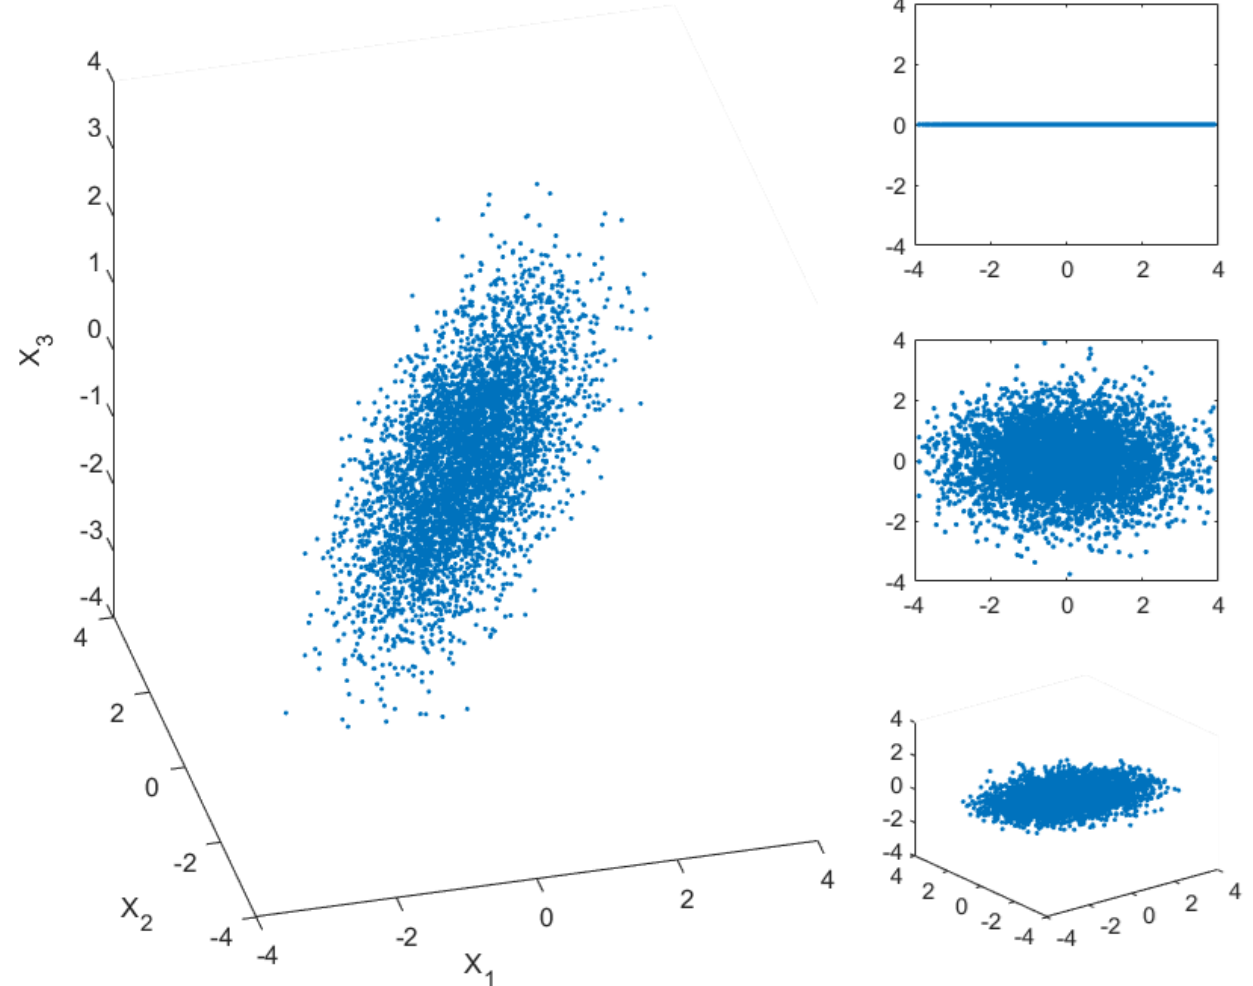


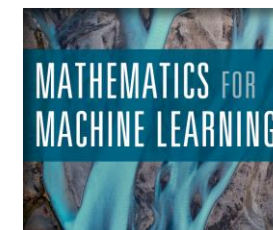
p 115



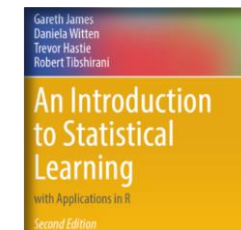
p 498

- **Principal Component Analysis (PCA):** Use the *eigenvectors*  $\mathbf{q}_1, \mathbf{q}_2, \mathbf{q}_3 \dots$  of the *covariance matrix*  $\mathbf{X}^T \mathbf{X}$  corresponding to the largest *eigenvalues*  $\lambda_1, \lambda_2, \lambda_3 \dots$





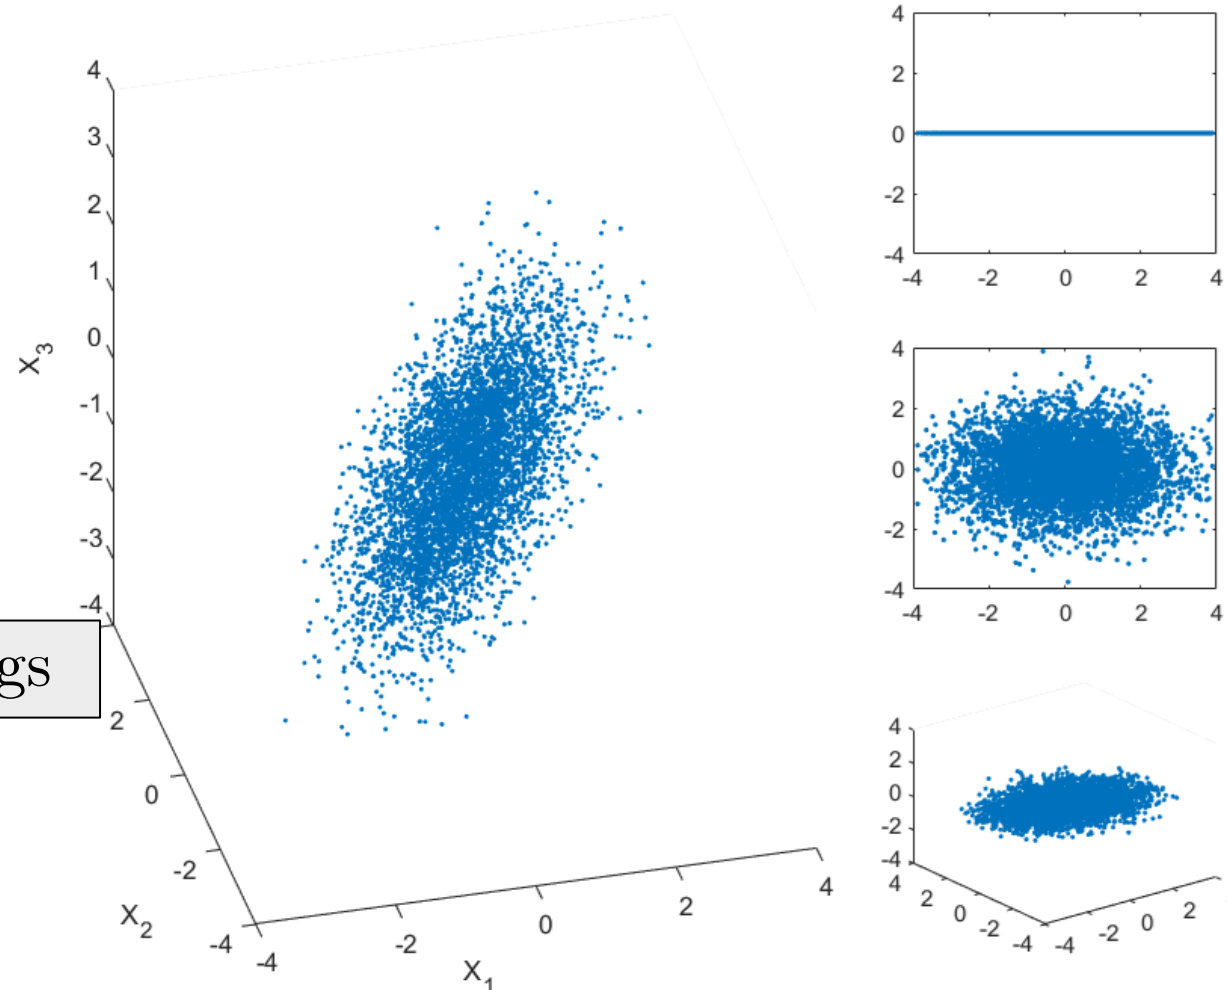
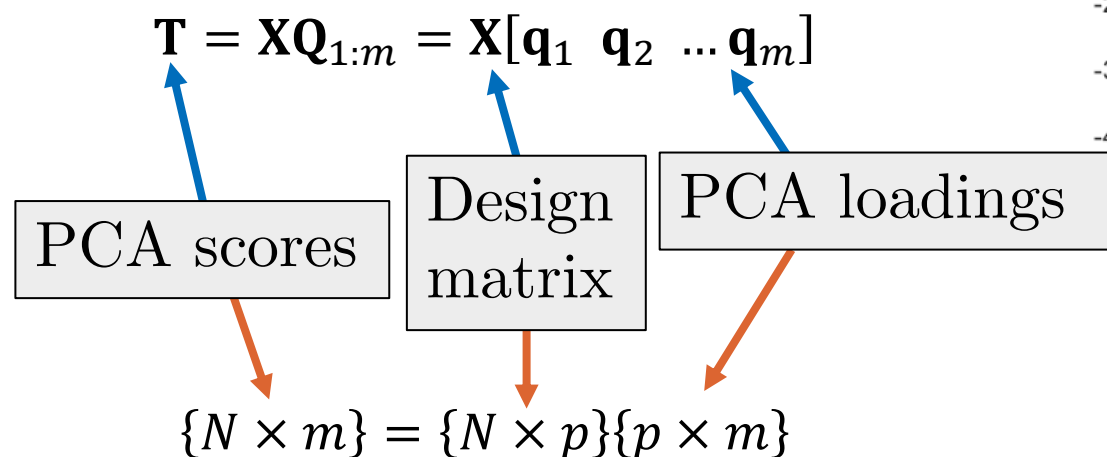
p 115



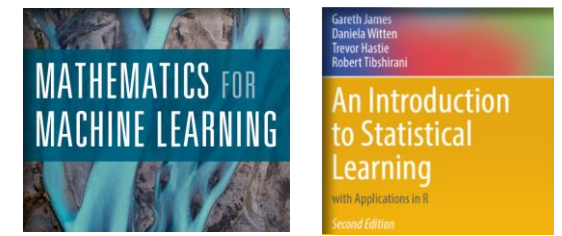
p 498

# Orthonormal matrix as “rotation”

- **Principal Component Analysis (PCA):** Use the *eigenvectors*  $\mathbf{q}_1, \mathbf{q}_2, \mathbf{q}_3 \dots$  of the *covariance matrix*  $\mathbf{X}^T \mathbf{X}$  corresponding to the largest *eigenvalues*  $\lambda_1, \lambda_2, \lambda_3 \dots$







p 115

p 498

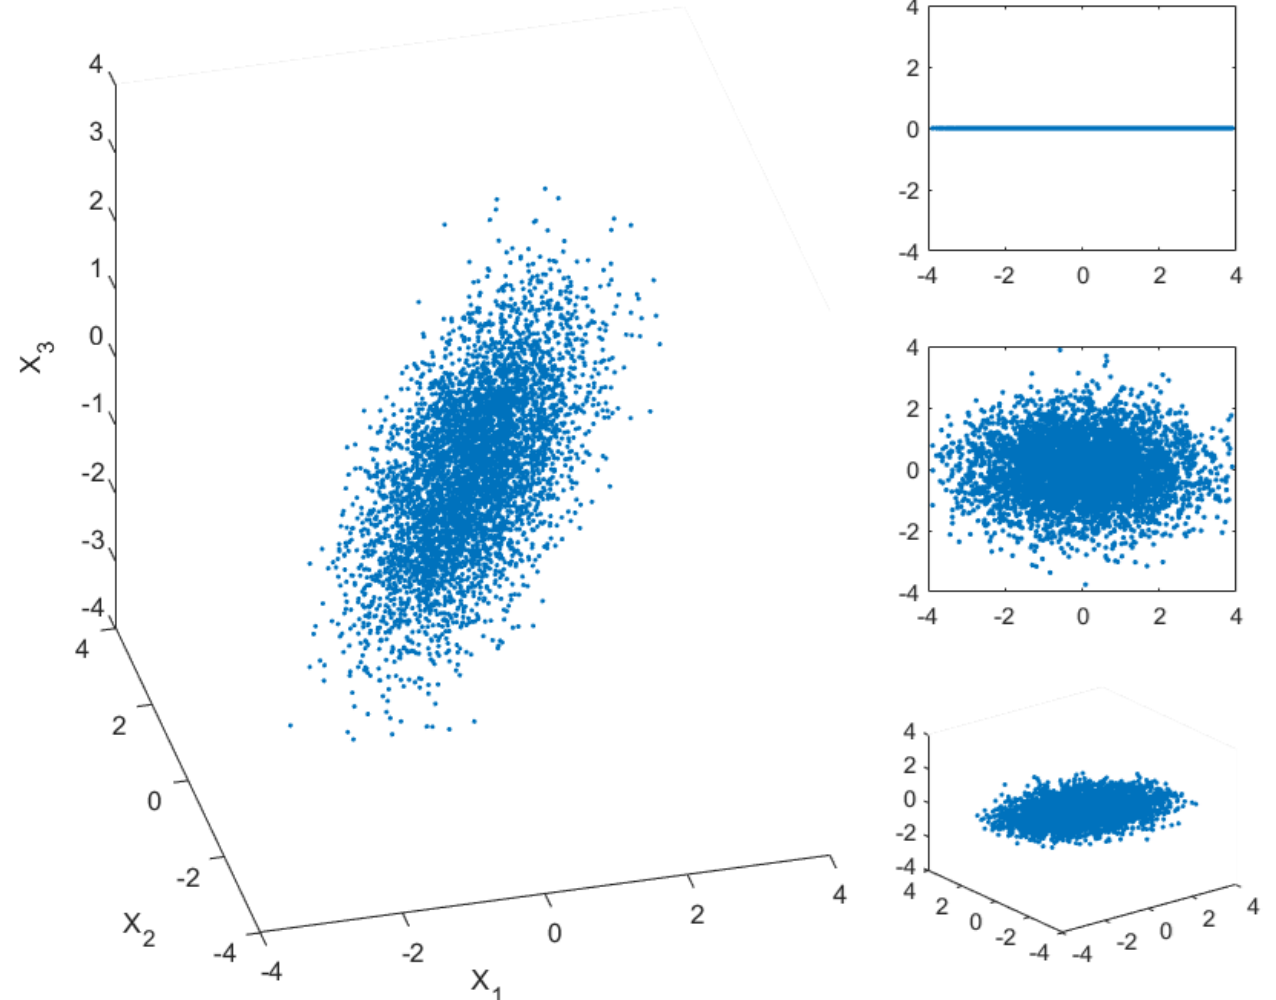
# Orthonormal matrix as “rotation”

- **Principal Component Analysis (PCA)**: Use the *eigenvectors*  $\mathbf{q}_1, \mathbf{q}_2, \mathbf{q}_3 \dots$  of the *covariance matrix*  $\mathbf{X}^T \mathbf{X}$  corresponding to the largest *eigenvalues*  $\lambda_1, \lambda_2, \lambda_3 \dots$

$$\mathbf{T} = \mathbf{X}\mathbf{Q}_{1:m} = \mathbf{X}[\mathbf{q}_1 \ \mathbf{q}_2 \ \dots \ \mathbf{q}_m]$$

- Because the covariance matrix  $\mathbf{X}^T \mathbf{X}$  is symmetric, the eigenvalues are orthonormal:

$$\mathbf{X}\mathbf{Q} \rightarrow \text{rotation}$$



# Orthonormal matrix as “rotation”

$$\mathbf{T} = \mathbf{X}[\mathbf{q}_1 \ \mathbf{q}_2 \ \dots \ \mathbf{q}_m]$$

$$\begin{bmatrix} \vdots & \vdots \\ t_{i-1,1} & t_{i-1,2} \\ t_{i,1} & t_{i,2} \\ t_{i+1,1} & t_{i+1,1} \\ \vdots & \vdots \end{bmatrix} = \begin{bmatrix} \vdots & \vdots & \vdots \\ x_{i-1,1} & x_{i-1,2} & x_{i-1,3} \\ x_{i,1} & x_{i,2} & x_{i,3} \\ x_{i+1,1} & x_{i+1,2} & x_{i+1,3} \\ \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} q_{1,1} & q_{2,1} \\ q_{1,2} & q_{2,2} \\ q_{1,3} & q_{2,3} \end{bmatrix}$$

# Orthonormal matrix as “rotation”

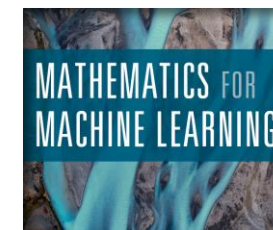
$$\mathbf{T} = \mathbf{X}[\mathbf{q}_1 \ \mathbf{q}_2 \ \dots \ \mathbf{q}_m]$$

$$\begin{bmatrix} \vdots & \vdots \\ t_{i-1,1} & t_{i-1,2} \\ t_{i,1} & t_{i,2} \\ t_{i+1,1} & t_{i+1,1} \\ \vdots & \vdots \end{bmatrix} = \begin{bmatrix} \vdots & \vdots & \vdots \\ x_{i-1,1} & x_{i-1,2} & x_{i-1,3} \\ x_{i,1} & x_{i,2} & x_{i,3} \\ x_{i+1,1} & x_{i+1,2} & x_{i+1,3} \\ \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} q_{1,1} & q_{2,1} \\ q_{1,2} & q_{2,2} \\ q_{1,3} & q_{2,3} \end{bmatrix}$$

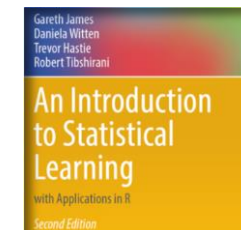
# Orthonormal matrix as “rotation”

$$\mathbf{T} = \mathbf{X}[\mathbf{q}_1 \ \mathbf{q}_2 \ \dots \ \mathbf{q}_m]$$

$$\begin{bmatrix} \vdots & \vdots \\ t_{i-1,1} & t_{i-1,2} \\ t_{i,1} & t_{i,2} \\ t_{i+1,1} & t_{i+1,2} \\ \vdots & \vdots \end{bmatrix} = \begin{bmatrix} \vdots & \vdots & \vdots \\ x_{i-1,1} & x_{i-1,2} & x_{i-1,3} \\ x_{i,1} & x_{i,2} & x_{i,3} \\ x_{i+1,1} & x_{i+1,2} & x_{i+1,3} \\ \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} q_{1,1} & q_{2,1} \\ q_{1,2} & q_{2,2} \\ q_{1,3} & q_{2,3} \end{bmatrix}$$



p 317



p 251

# Principal component regression

1. Project data to a lower dimensional space using the first  $m < p$  principal component loadings,  $\mathbf{Q}_{1:m}$ , with  $\mathbf{q}_i$  ordered according to  $\lambda_1 > \lambda_2 > \lambda_3 \dots$
2. Regress response variables onto  $m$  principal component scores

$$\mathbf{y} = \mathbf{XQ}_{1:m}\boldsymbol{\beta}$$

- Correlated variables are combined into single features
- Fewer predictors, fewer parameters, decrease model variance

```
%% Example 11: Use PCA regression to predict time series
```

```
[loadings, ~, ~, ~, explained] = pca(X, 'NumComponents',20);
```

```
% Plot the variance explained...
```

```
clf
```

```
subplot(2,1,1)
```

```
bar(explained);
```

```
⋮
```

```
% Fit the linear model to the reduced
```

```
% set of predictors  $X \cdot Q$ 
```

```
PCA_mdl.Q = loadings(:, 1:4);
```

```
T = X*PCA_mdl.Q;
```

```
mdl = fitlm(T, y, 'Intercept', false);
```

```
PCA_mdl.beta = mdl.Coefficients.Estimate;
```

```
% Simulate the model response
```

```
y_PCA = PredictTimeSeries(PCA_mdl, Data, L);
```

```
% Plot and compare the results
```

```
⋮
```

```
%% Example 11: Use PCA regression to predict time series
```

```
[loadings, ~, ~, ~, explained] = pca(X, 'NumComponents', 20);
```

```
% Plot the variance explained...
```

```
clf
```

```
subplot(2,1,1)
```

```
bar(explained);
```

```
:
```

```
% Fit the linear model to the reduced
```

```
% set of predictors X*Q
```

```
PCA_mdl.Q = loadings(:, 1:4);
```

```
T = X*PCA_mdl.Q;
```

```
mdl = fitlm(T, y, 'Intercept', false);
```

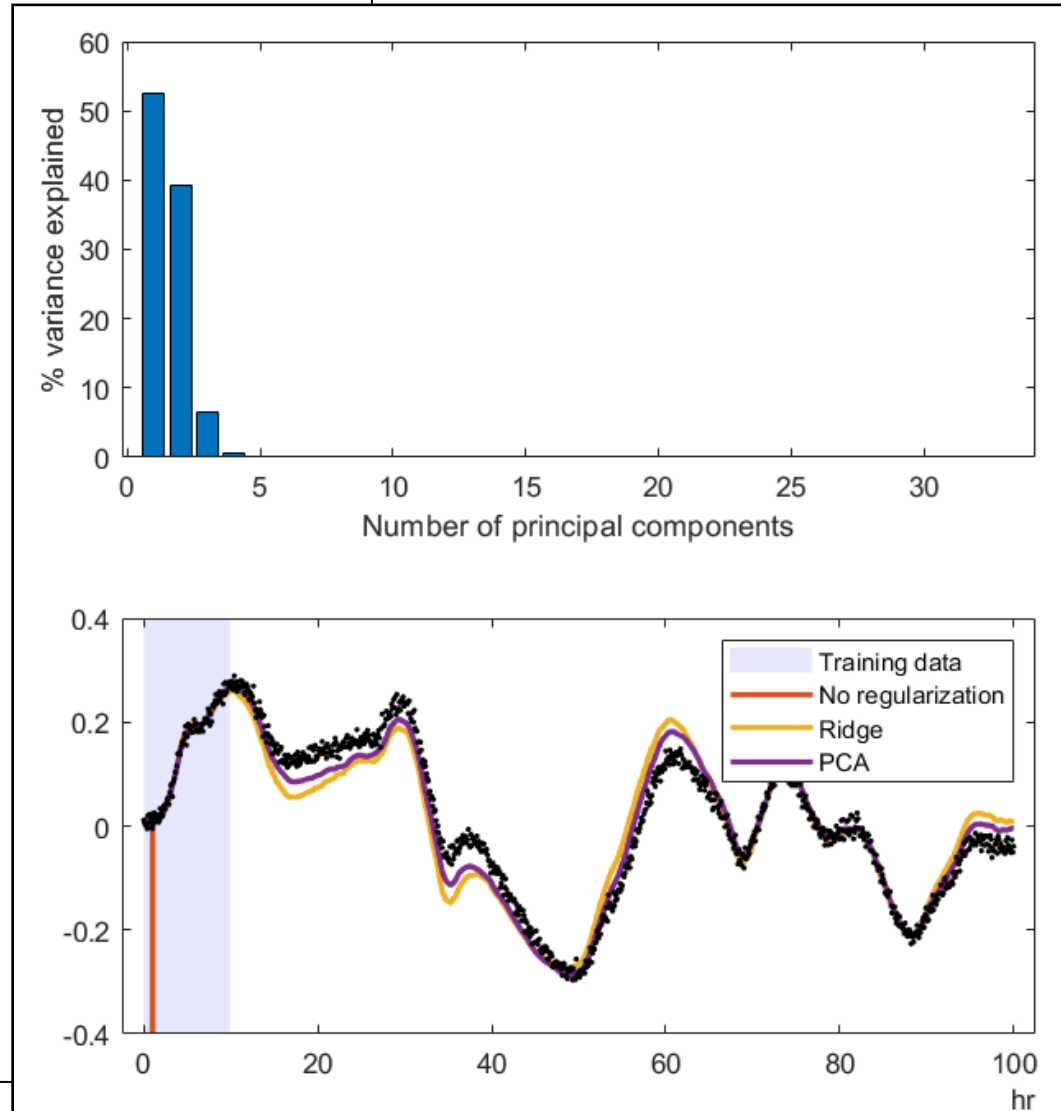
```
PCA_mdl.beta = mdl.Coefficients.Estimate;
```

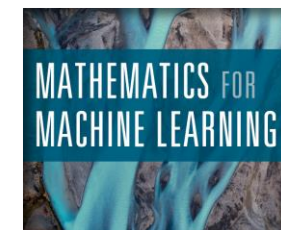
```
% Simulate the model response
```

```
y_PCA = PredictTimeSeries(PCA_mdl, Data, L);
```

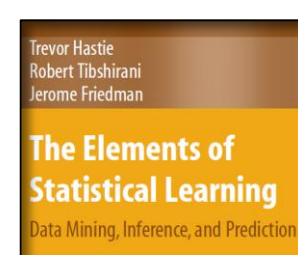
```
% Plot and compare the results
```

```
:
```





p 317



p 79

# Principal component regression

1. Project data to a lower dimensional space using the first  $m < p$  principal component loadings,  $\mathbf{Q}_{1:m}$ , with  $\mathbf{q}_i$  ordered according to  $\lambda_1 > \lambda_2 > \lambda_3 \dots$
2. Regress response variables onto  $m$  principal component scores

$$\mathbf{y} = \mathbf{XQ}_{1:m}\boldsymbol{\beta}$$

- Correlated variables are combined into single features
- Fewer predictors, fewer parameters, decrease model variance



# Principal component regression

Find components  $\mathbf{q}_i$  such that:

$$\mathbf{q}_i = \arg \max_{\mathbf{v}} \{\text{Var}(\mathbf{X}\mathbf{v})\} = \arg \max_{\mathbf{v}} \{(\mathbf{X}\mathbf{v}) \cdot (\mathbf{X}\mathbf{v})\},$$

$$\text{s.t. } |\mathbf{v}| = 1 \quad (\mathbf{X}\mathbf{v}) \cdot (\mathbf{X}\mathbf{q}_j) = \mathbf{0} \quad \forall j < i$$

The input is selected to maximize  
variance in the input

The response  $\mathbf{y}$  is not considered in  
constructing the input directions

# Partial least squares regression

Find components  $\mathbf{q}_i$  such that:

$$\begin{aligned}\mathbf{q}_i &= \arg \max_{\mathbf{v}} \{ \text{Corr}^2(\mathbf{y}, \mathbf{X}\mathbf{v}) \text{Var}(\mathbf{X}\mathbf{v}) \} \\ &= \arg \max_{\mathbf{v}} \{ (\mathbf{y}) \cdot (\mathbf{X}\mathbf{v}) \times (\mathbf{X}\mathbf{v}) \cdot (\mathbf{X}\mathbf{v}) \},\end{aligned}$$

$$\text{s.t. } |\mathbf{v}| = 1 \quad (\mathbf{X}\mathbf{v}) \cdot (\mathbf{X}\mathbf{q}_j) = \mathbf{0} \quad \forall j < i$$

The input is selected to maximize correlation with the response  $\mathbf{y}$  as well as variance in the input

# Partial least squares regression

1. Set  $i = 0$

2. Let  $\tilde{\mathbf{X}}^{(0)} = \mathbf{X}$

3. Let  $i \leftarrow i + 1$

4. Obtain loading direction:

$$\mathbf{q}_i = \sum_{j=1}^p \left( \tilde{\mathbf{X}}_j^{(i-1)} \cdot \mathbf{y} \right) \tilde{\mathbf{X}}_j^{(i-1)}$$

5. Regress coefficient:

$$\beta_i \leftarrow \mathbf{y} = \tilde{\mathbf{X}}^{(i-1)} \mathbf{q}_i \beta_i$$

6. Generate design matrix orthogonal to  $\mathbf{q}_i$ :

$$\tilde{\mathbf{X}}_j^{(i)} = \tilde{\mathbf{X}}_j^{(i-1)} - \left( \frac{\mathbf{q}_i \cdot \tilde{\mathbf{X}}_j^{(i-1)}}{\mathbf{q}_i \cdot \mathbf{q}_i} \right) \mathbf{q}_i$$

7. Repeat steps 3-7 until  $i = m$

# Partial least squares regression

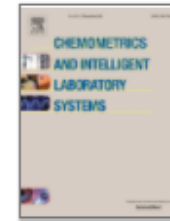
- MATLAB implementation relies on slightly different algorithm (SIMPLS)

[https://doi.org/10.1016/0169-7439\(93\)85002-X](https://doi.org/10.1016/0169-7439(93)85002-X)



Chemometrics and Intelligent Laboratory Systems

Volume 18, Issue 3, March 1993, Pages 251-263



---

## SIMPLS: An alternative approach to partial least squares regression

Sijmen de Jong 

```
%% Example 12: Use PLS regression to predict time series
```

```
[loadings, ~, ~, ~, ~, explained] = plsregress(X, y, 4);
```

```
% Plot the variance explained..
```

```
?
```

```
?
```

```
?
```

```
% Fit the linear model...
```

```
?
```

```
?
```

```
?
```

```
?
```

```
% Simulate the model response
```

```
?
```

```
% Plot and compare the results
```

```
⋮
```

```
%% Example 12: Use PLS regression to predict time series
```

```
[loadings, ~, ~, ~, ~, explained] = plsregress(X, y, 4);
```

```
% Plot the variance explained..
```

```
clf
```

```
subplot(2,1,1)
```

```
bar(explained);
```

```
% Fit the linear model...
```

```
PLS_md1.Q = loadings(:, 1:4);
```

```
T = X*PLS_md1.Q;
```

```
mdl = fitlm(T, y, 'Intercept', false);
```

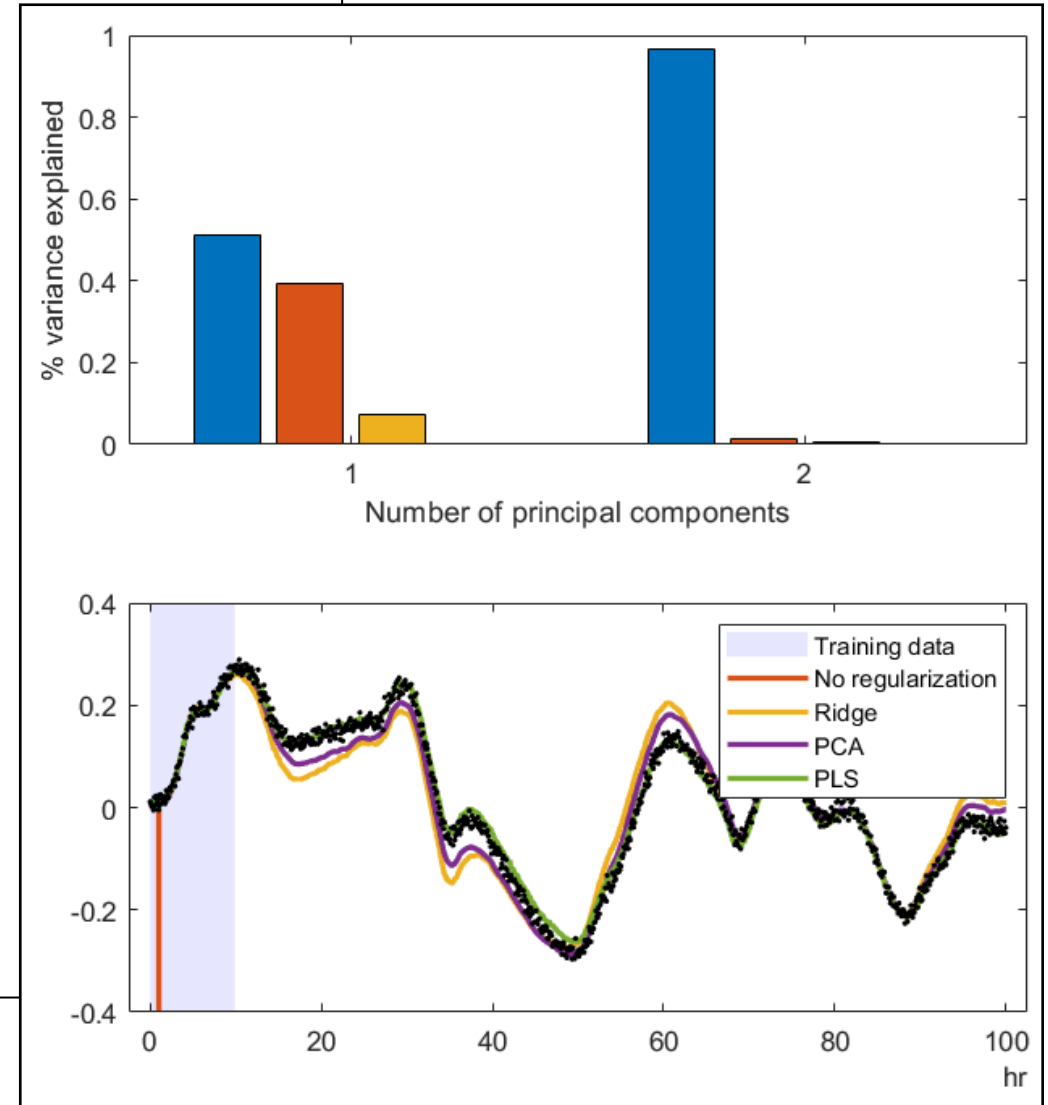
```
PLS_md1.beta = mdl.Coefficients.Estimate;
```

```
% Simulate the model response
```

```
y_PLS = PredictTimeSeries(PLS_md1, Data, L);
```

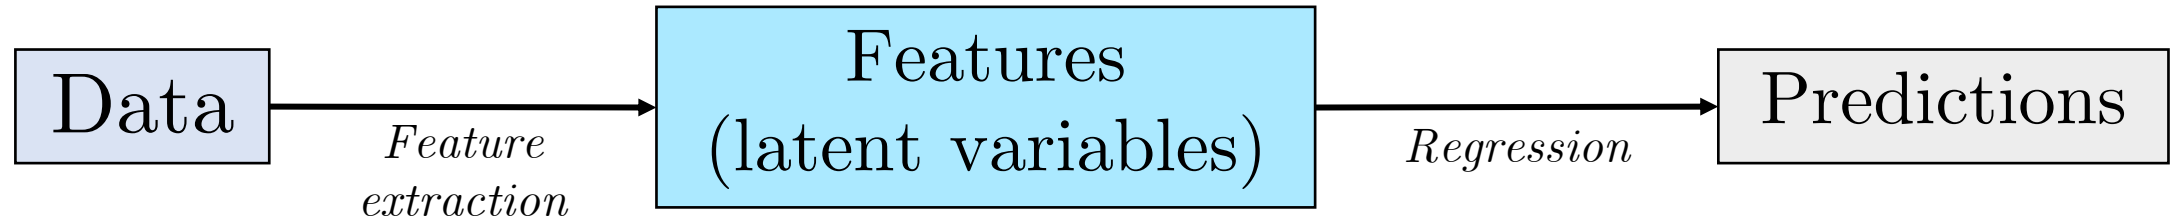
```
% Plot and compare the results
```

```
:
```



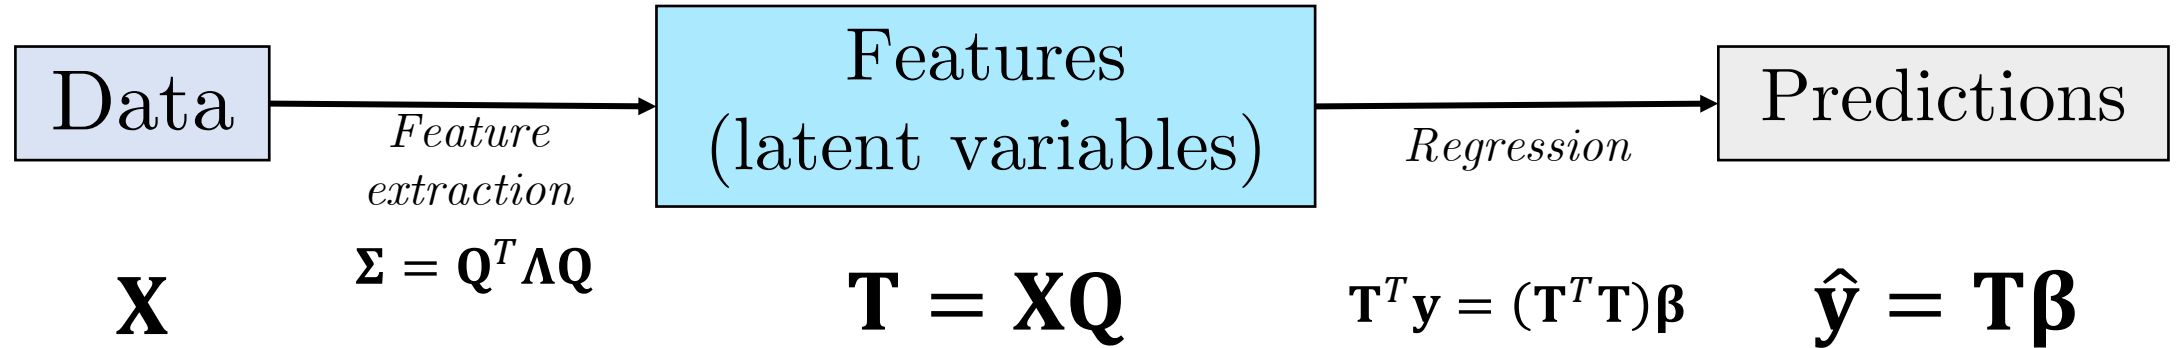
# Recap

- Model variance can be reduced by extracting the most important features, and discarding those that do not contribute



# Recap

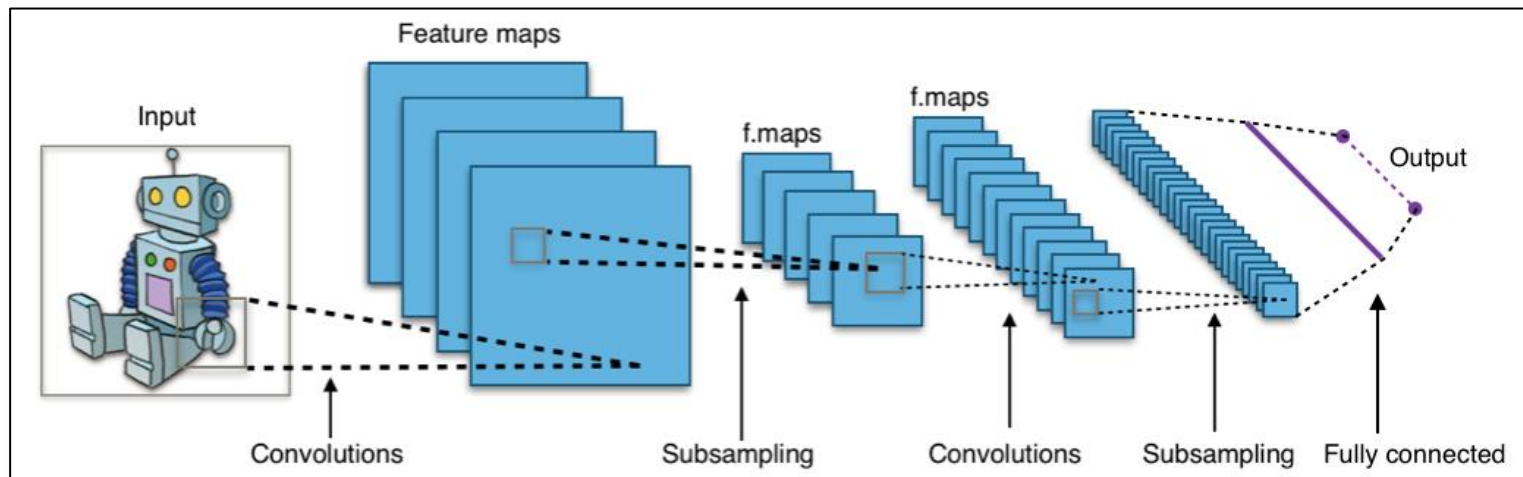
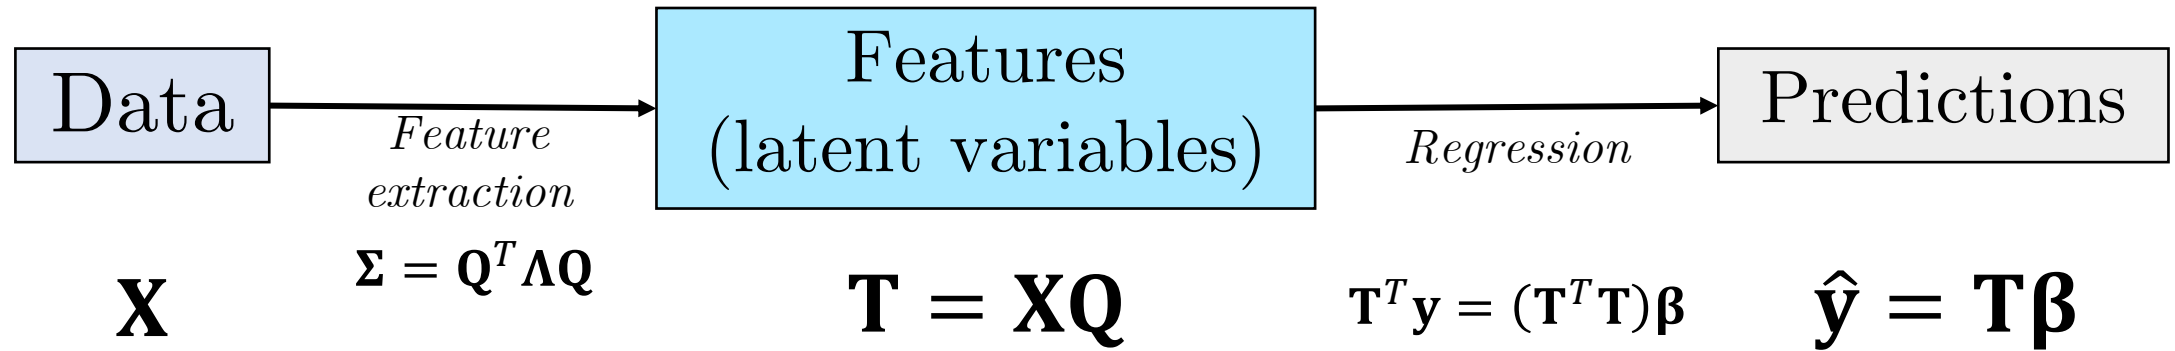
- Model variance can be reduced by extracting the most important features, and discarding those that do not contribute





# Recap

- Model variance can be reduced by extracting the most important features, and discarding those that do not contribute



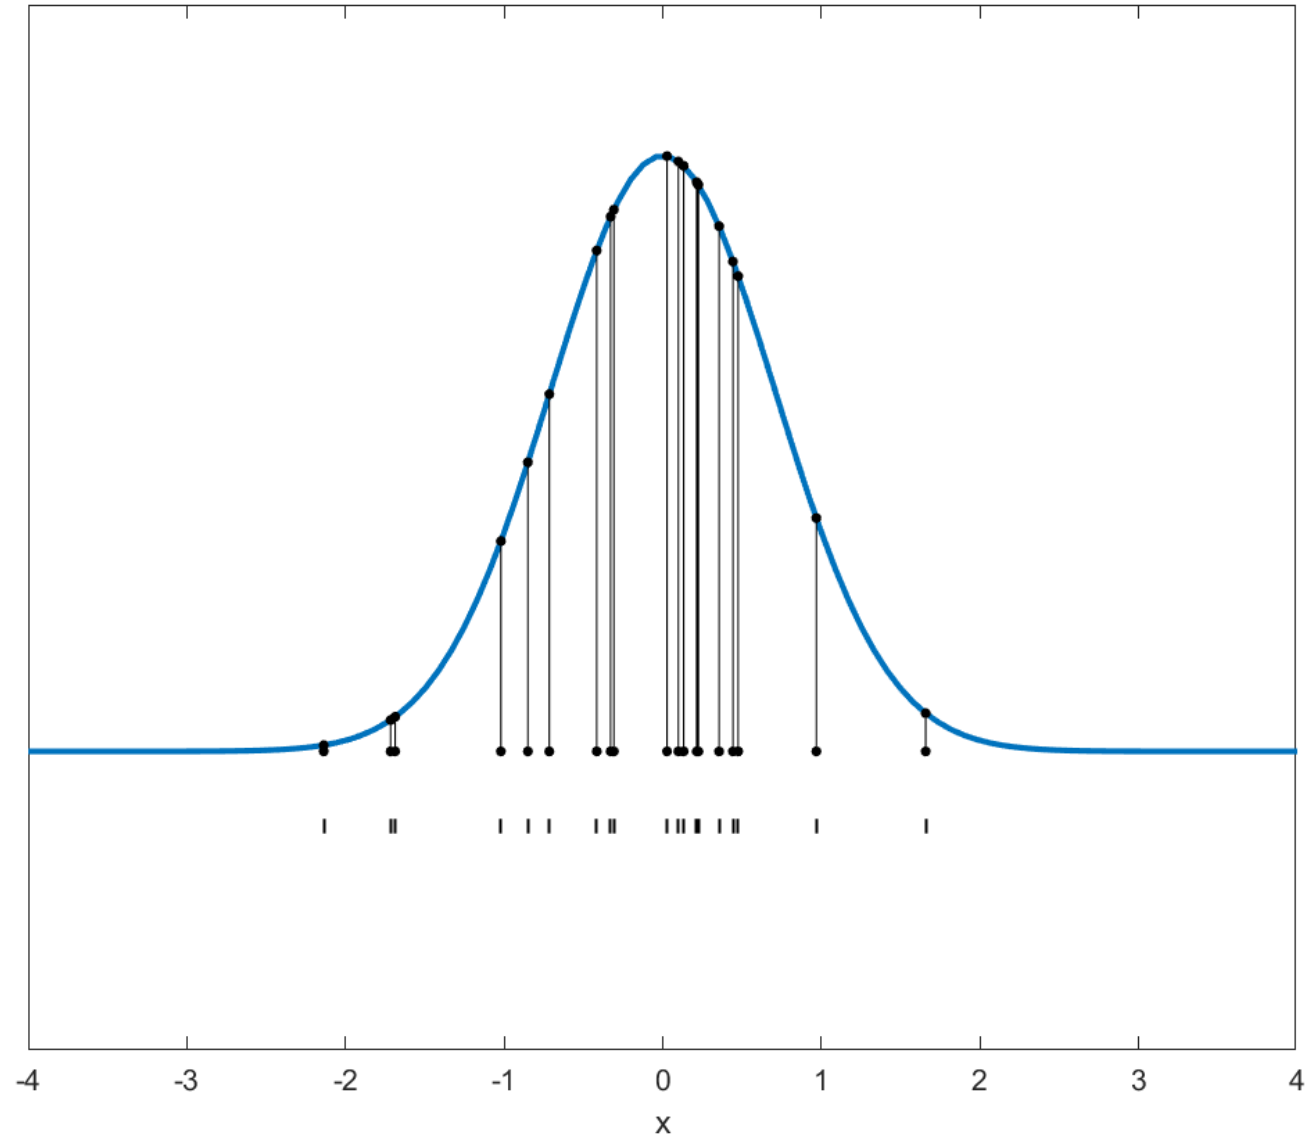
# Recap

- Feature extraction provides an alternative approach to decrease model variance and improve prediction accuracy
- PCA and PLS are examples of linear feature extraction
- Common in system identification: Canonical Correlation Analysis (CCA)
- Non-linear feature extraction / latent variable modelling is at the heart of many modern machine learning methods (deep learning, GP-LVM, etc)

# Why are the eigenvectors of the covariance matrix also the principal components?

- Consider data distributed normally with zero mean and unit variance
- The probability density of an observation  $x_i$  is given by:

$$p(x_i) = \frac{1}{\sqrt{2\pi}} \exp(-x_i^2)$$



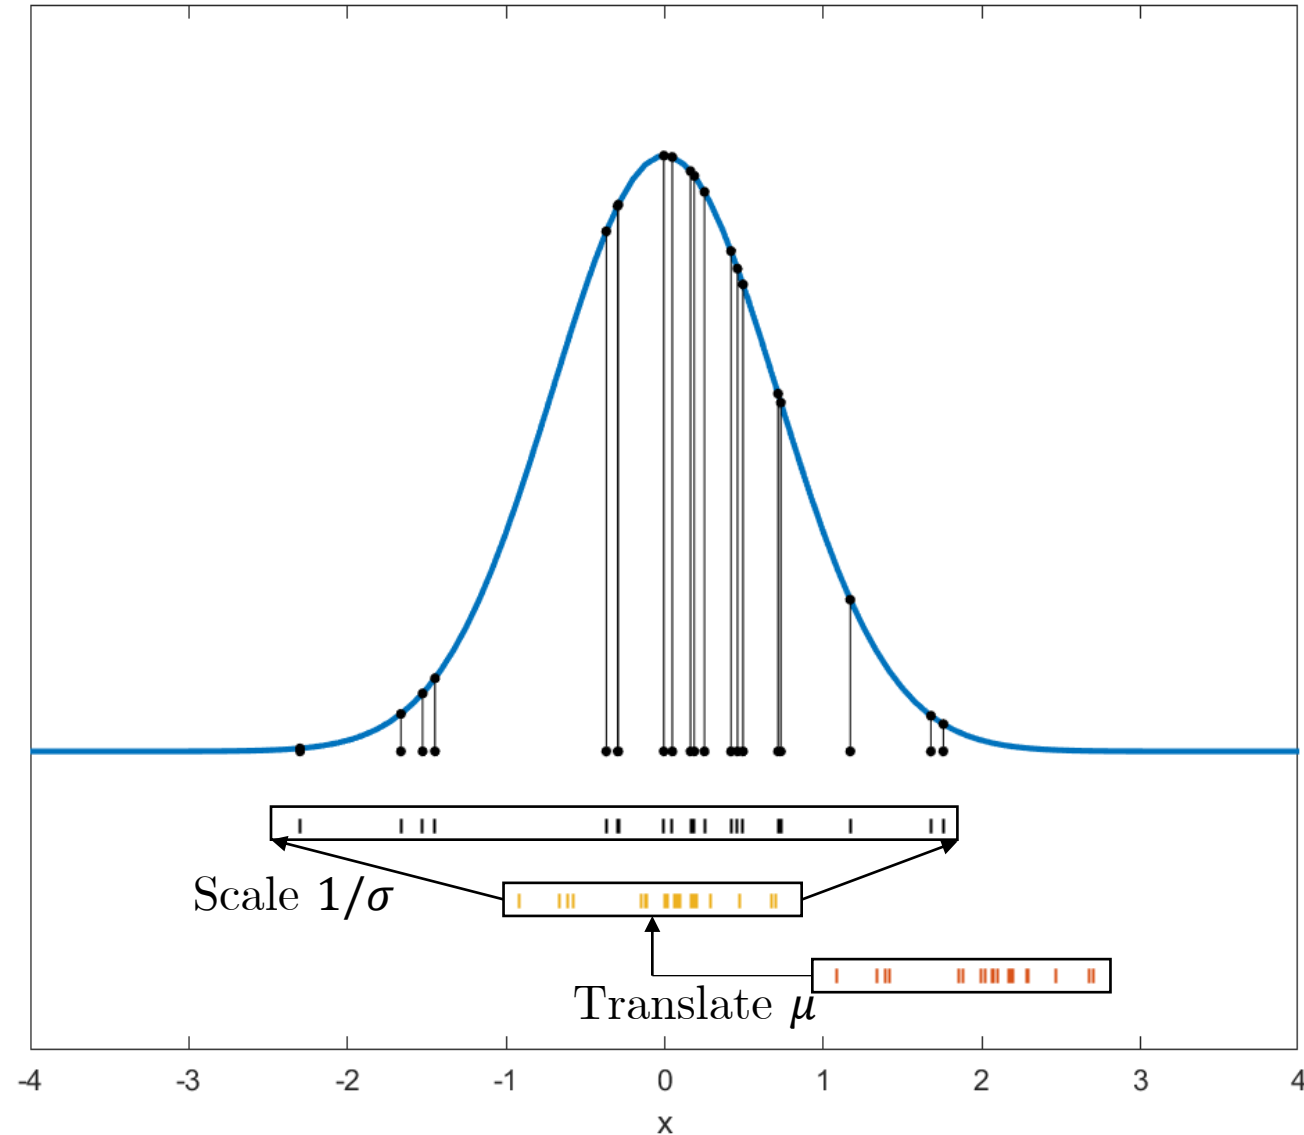
# Why are the eigenvectors of the covariance matrix also the principal components?

- Consider data distributed with mean  $\mu$  and variance  $\sigma^2$
- The probability density is found by translating each  $x_i$  by  $\mu$ , then scaling by  $\sigma$

$$\tilde{x}_i = \frac{x_i - \mu}{\sigma}$$

- The probability density of an observation  $x_i$  is given by:

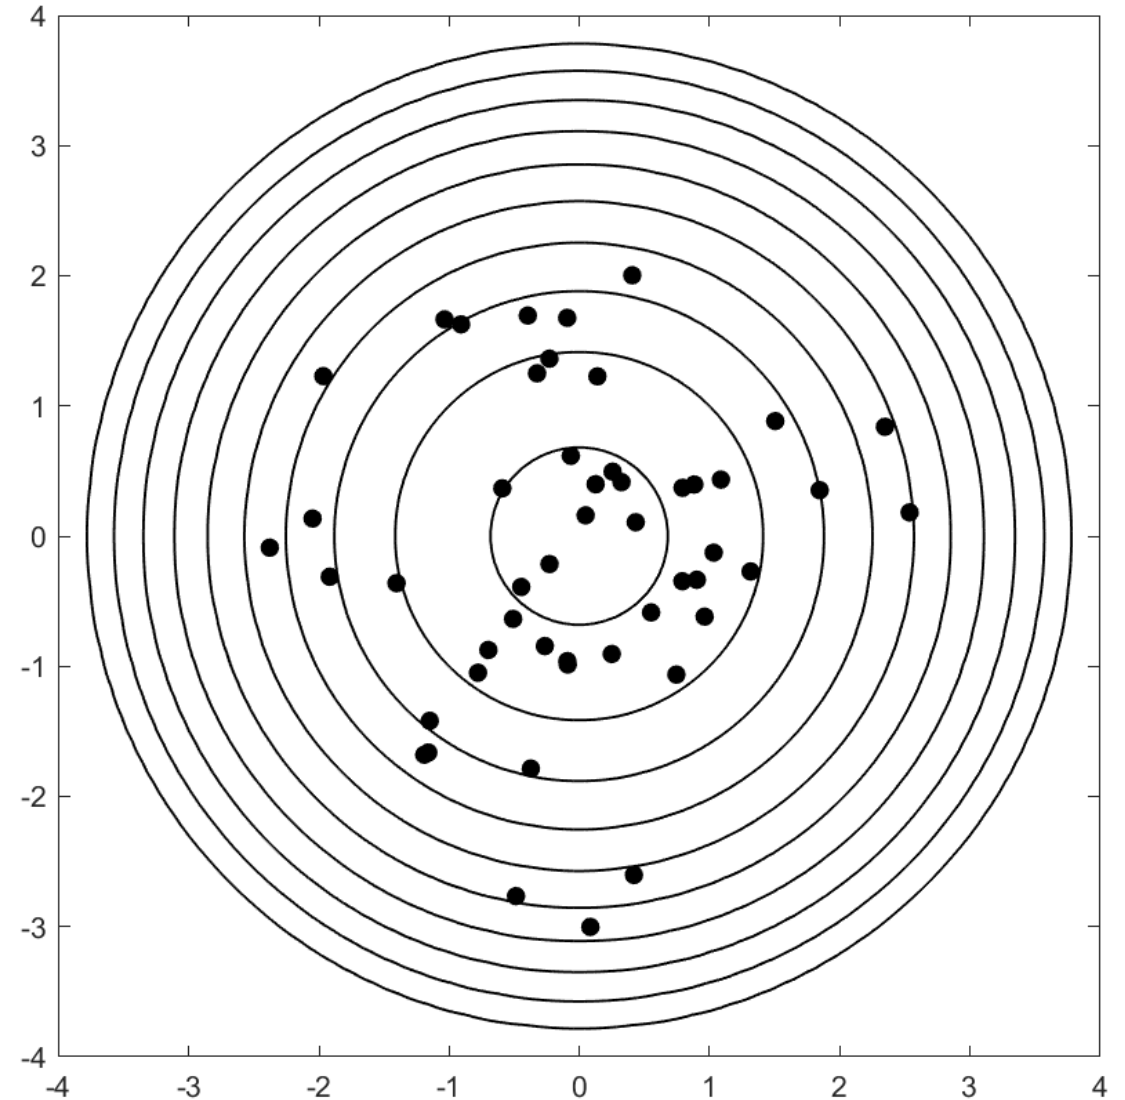
$$p(x_i) = \frac{1}{\sqrt{2\pi}} \exp(-\tilde{x}_i^2)$$



# Why are the eigenvectors of the covariance matrix also the principal components?

- Consider 2-dimensional data distributed normally with zero mean and unit variance
- The probability density of an observation  $(x_{i,1}, x_{i,2})$  is given by:

$$\begin{aligned} p(\mathbf{x}_i) &= \frac{1}{2\pi} \exp\left(-\left(x_{i,1}^2 + x_{i,2}^2\right)\right) \\ &= \frac{1}{2\pi} \exp\left(-\begin{bmatrix} x_{i,1} & x_{i,2} \end{bmatrix} \begin{bmatrix} x_{i,1} \\ x_{i,2} \end{bmatrix}\right) \\ &= \frac{1}{2\pi} \exp(-\mathbf{x}_i \mathbf{x}_i^T) \end{aligned}$$



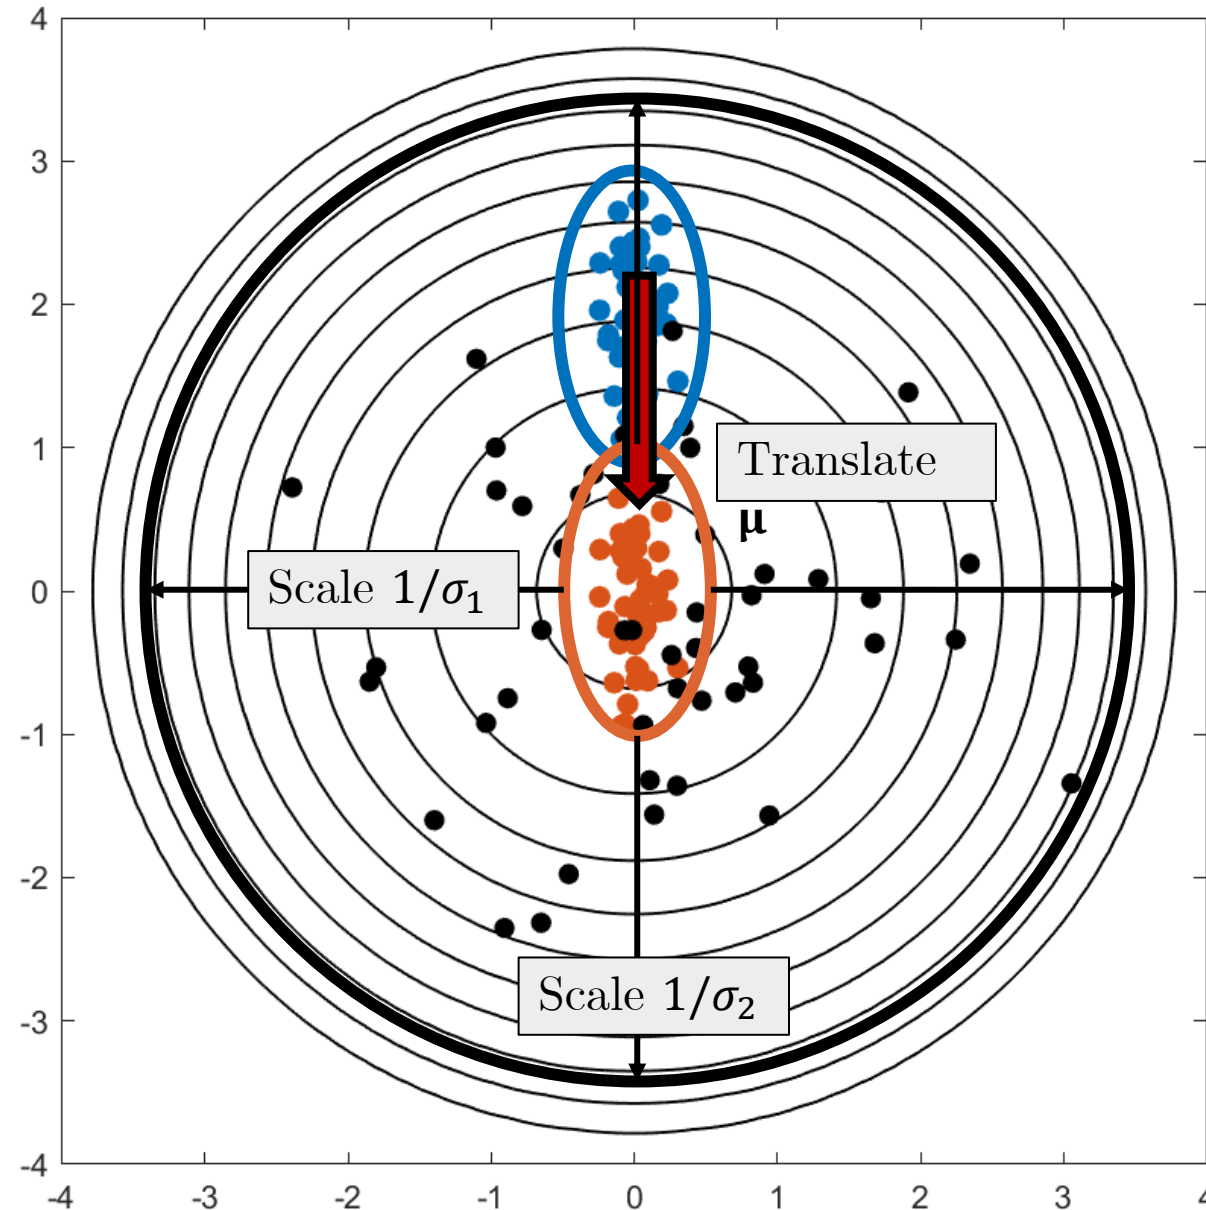
# Why are the eigenvectors of the covariance matrix also the principal components?

- Consider 2-dimensional data distributed normally with mean  $\boldsymbol{\mu}$  and covariance  $\Sigma = \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{bmatrix}$
- The data is translated and scaled by

$$\begin{aligned} [\tilde{x}_{i,1} \quad \tilde{x}_{i,2}] &= \begin{bmatrix} \frac{x_{i,1} - \mu_1}{\sigma_1} & \frac{x_{i,2} - \mu_2}{\sigma_2} \end{bmatrix} \\ &= (\mathbf{x}_i - \boldsymbol{\mu}) \begin{bmatrix} 1/\sigma_1 & 0 \\ 0 & 1/\sigma_2 \end{bmatrix} \\ &= (\mathbf{x}_i - \boldsymbol{\mu}) \Sigma^{-1/2} \end{aligned}$$

- The probability density of an observation  $(x_{i,1}, x_{i,2})$  is given by:

$$p(\mathbf{x}_i) = \frac{1}{2\pi} \exp(-\tilde{\mathbf{x}}_i \tilde{\mathbf{x}}_i^T)$$



# Why are the eigenvectors of the covariance matrix also the principal components?

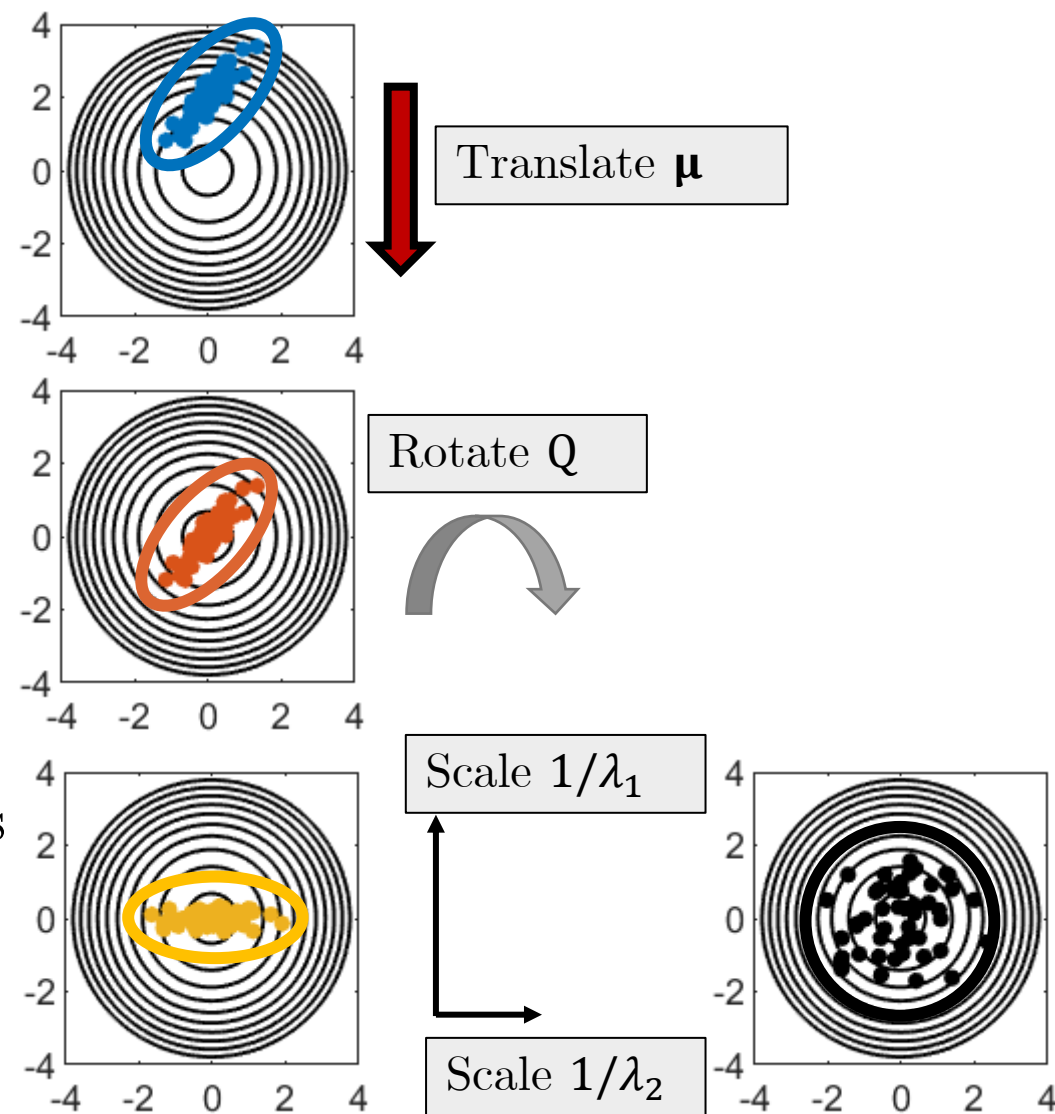
- Consider 2-dimensional data distributed normally with mean  $\boldsymbol{\mu}$  and covariance  $\Sigma = \begin{bmatrix} \sigma_1^2 & \sigma_{12}^2 \\ \sigma_{12}^2 & \sigma_2^2 \end{bmatrix}$

- The data is translated, rotated, and scaled by

$$\begin{aligned} [\tilde{x}_{i,1} \quad \tilde{x}_{i,2}] &= (\mathbf{x}_i - \boldsymbol{\mu}) \begin{bmatrix} q_{11} & q_{21} \\ q_{12} & q_{22} \end{bmatrix} \begin{bmatrix} 1/\Lambda_1 & 0 \\ 0 & 1/\Lambda_1 \end{bmatrix} \\ &= (\mathbf{x}_i - \boldsymbol{\mu}) \mathbf{Q} \Lambda^{-1/2} = (\mathbf{x}_i - \boldsymbol{\mu}) \Sigma^{-1/2} \end{aligned}$$

- The probability density of an observation  $(x_{i,1}, x_{i,2})$  is given by:

$$p(\mathbf{x}_i) = \frac{1}{2\pi} \exp(-\tilde{\mathbf{x}}_i \tilde{\mathbf{x}}_i^T)$$



## Why are the eigenvectors of the covariance matrix also the principal components?

- The eigendecomposition of the covariance matrix yields:

$$\Sigma = Q\Lambda Q^T = (Q\Lambda^{1/2})(Q\Lambda^{1/2})^T$$

$$\Sigma^{-1} = Q\Lambda Q^T = (Q\Lambda^{-1/2})(Q\Lambda^{-1/2})^T$$

$$p(\mathbf{x}_i) = \frac{1}{2\pi\sqrt{|\Sigma|}} \exp\left(-\frac{1}{2}(\mathbf{x}_i - \boldsymbol{\mu})\Sigma^{-1}(\mathbf{x}_i - \boldsymbol{\mu})^T\right)$$

$$= \frac{1}{2\pi\sqrt{|\Sigma|}} \exp\left(-\frac{1}{2}\left[(\mathbf{x}_i - \boldsymbol{\mu})(Q\Lambda^{-1/2})\right]\left[(\mathbf{x}_i - \boldsymbol{\mu})(Q\Lambda^{-1/2})\right]^T\right)$$